

Data in/and \mathbb{R}

EC 425/525, Lab 2

Edward Rubin

12 April 2019

Prologue

Schedule

Last time

Getting to know R—objects, functions, *etc.*

Today

Working with data in R.

- The `data.frame` class
- The `dplyr` package

Upcoming

Due Monday Step 1 of our research-project proposal.

Matrices

Quick review

1. `mat ← matrix(data = 1:10, ncol = 2)` creates a 5×2 `matrix` object containing the numbers 1 through 10 (filled by column).
2. `mat[1,]` grabs the first row of our matrix `mat`.
3. `mat[3,2] ← NA` assigns `NA` to row-3 column-2 element of `mat`.
4. `head(mat, 3)` returns up to the first three rows of `mat`.
5. `matrix(data = rnorm(100), ncol = 10)` creates a 10×10 matrix filled with random draws from $N(\mu = 0, \sigma^2 = 1)$.
6. `mat[3,2] ← "Carrots"` assigns the `character` object "Carrots" to the `[3,2]` element of `mat`, forcing all elements of `mat` to `character`.

Matrices

Next steps

Matrices are convenient two-dimensional arrays on which math "works."[†]

But matrices also require all elements to be of the same class.

Q What if we have datasets whose variables (columns) have different classes?

A We need a more flexible table-like object for our data.

Maybe a `data.table`? Or a `data.frame`?

We'll start with `data.frame`.

We will spend a good amount of time on data frames, as they make up a huge part of your workflow.

[†] At least for `numeric` and `logical` matrices.

Data frames

A `data.frame` is R's base, spreadsheet-like object that holds variables.

Example

```
#>      id first_name fave_num is_tired loves_econ
#> 1     1      Karmin      68      TRUE      FALSE
#> 2     2   Raychelle      57      TRUE       TRUE
#> 3     3    Jemelle      10      TRUE       TRUE
#> 4     4      Yusif      90      TRUE       TRUE
#> 5     5  Catherine      24      TRUE       TRUE
#> 6     6      Glory       4      TRUE       TRUE
#> 7     7    Kaelah      33     FALSE       TRUE
#> 8     8    Lysette      96      TRUE       TRUE
#> 9     9     Cisco      89      TRUE       TRUE
#> 10    10    Harman      69      TRUE       TRUE
#> 11    11  Jennelle      64      TRUE       TRUE
#> 12    12   Crayton     100      TRUE       TRUE
```

Data frames

A `data.frame` is R's base, spreadsheet-like object that holds variables.

Example

```
#>           name height mass gender homeworld species
#> 1   Luke Skywalker   172   77   male   Tatooine   Human
#> 2         C-3PO    167   75  <NA>   Tatooine   Droid
#> 3         R2-D2     96   32  <NA>     Naboo   Droid
#> 4   Darth Vader    202  136   male   Tatooine   Human
#> 5   Leia Organa    150   49 female Alderaan   Human
#> 6     Owen Lars    178  120   male   Tatooine   Human
#> 7 Beru Whitesun lars   165   75 female   Tatooine   Human
#> 8         R5-D4     97   32  <NA>   Tatooine   Droid
#> 9 Biggs Darklighter   183   84   male   Tatooine   Human
#> 10   Obi-Wan Kenobi   182   77   male   Stewjon   Human
#> 11   Anakin Skywalker   188   84   male   Tatooine   Human
#> 12   Wilhuff Tarkin   180   NA   male     Eriadu   Human
```

Data frames

Creation

The `data.frame()` function creates... `data.frame` objects.

You'll generally define data frames by passing the function
(1) column names and (2) values for the columns.

```
data.frame(var1 = 1:5, var2 = "apple", var3 = rnorm(5))
```

You can also assign the values using already-existing objects, *e.g.*,

```
# An object with value  
tmp ← rnorm(5)  
# Creating the data frame  
data.frame(var1 = 1:5, var2 = "apple", var3 = tmp)
```


Data frames

Creation

```
# Creating the data frame
```

```
data.frame(var1 = 1:5, var2 = "apple", var3 = rnorm(5))
```

```
#>   var1 var2      var3
#> 1     1 apple -0.6250393
#> 2     2 apple -1.6866933
#> 3     3 apple  0.8377870
#> 4     4 apple  0.1533731
#> 5     5 apple -1.1381369
```

(What a beauty.)

Notice that R assumes we want to repeat "apple" for the entire column.

Data frames

Creation

You can also create data frames from other objects (e.g., matrices) using the function `as.data.frame()`[†].

However, your data frame's columns will only have names if your matrix's columns had names.

[†] Or just plain, old `data.frame()`.

Data frames

Indexing

Consider a data frame `our_df ← data.frame(x = 1:3, y = 4:6, z = 7:9)`.

Option 1 Index data frames just as you index matrices in R.

- `our_df[1,1]` grabs the value in the first row of the first variable.
- `our_df[2,]` returns the second row of `our_df` (as a data frame).
- `our_df[,3]` returns the third column (variable) of `our_df` (as a vector).

Option 2 Reference values/variables using columns' names.

- `our_df$x` returns the column named `x` (as a vector). **New:** `$`
- `our_df[, "x"]` returns the column named `x` (as a vector).
- `our_df["x"]` returns the column named `x` (as a data frame).
- `our_df[, c("x", "y")]` returns a data frame with variables `"x"` and `"y"`.

Data frames

Names (of columns)

The columns (variables) in your data frame have names.[†]

Q What if you want to see/know those names?

A You've got a few options.

1. The `names()` function returns the *names* of an object.
1. `head(your_df)` will show you the first 6 rows of `your_df`.
Note: May provide too much output if you have a lot of columns.
1. In RStudio: `View(your_df)` or look in your Environment tab.

[†] If you don't name the columns, then R will.

Data frames

Naming

The `names()` function will also help you rename any/all variables.

Change the names of **all variables** (include a name for each variable):

```
# Set new names  
names(our_df) ← c("name1", "name2", "name3")
```

Change the name of **the second variable** (only):

```
# Set new names  
names(our_df)[2] ← "name2"
```

Data frames

Adding variables

Just as we referenced **existing** variables using `$var_name`, we can create **new** variables using `$new_var`, e.g.,

```
# Add a variable to our_df  
our_df$new_var ← 1:100
```

If you want to use existing columns to create a new variable

```
# Create interaction: xy = x * y  
our_df$xy ← our_df$x * our_df$y
```

Q Isn't there a better/faster/less-typing way?

A Yes. Enter `dplyr` (also: `data.table`, which we'll leave for the future).

dplyr

Intro

It's a package. `dplyr` is not installed by default, so you'll need to install it.[†]

`dplyr` is part of the `tidyverse` (Hadleyverse), and it follows a grammar-based approach to programming/data work.

- `data` compose the subjects of your stories
- `dplyr` provides the *verbs* (action words) :
`filter()`, `mutate()`, `select()`, `group_by()`, `summarize()`, `arrange()`

Bonus `dplyr` is pretty fast and able to interact with SQL databases.

[†] or just `install.packages("dplyr")` after loading `pacman`.

Manipulating variables: `mutate()`

`dplyr` streamlines adding/manipulating variables in your data frame.

Function `mutate(.data, ...)`

- **Required argument** `.data`, an existing data frame
- **Additional arguments** Names and values of the new variables
- **Output** An updated data frame

Example

```
mutate(.data = our_df, new1 = 7, new2 = x * y)
```


dplyr

mutate()

Example Take the data frame

```
my_df <- data.frame(x = 1:4, y = 5:8)
```

`mutate()` allows us to create many new variables with one call.

```
mutate(.data = my_df,  
  xy = x * y,  
  x2 = x^2,  
  y2 = y^2,  
  xy2 = xy^2,  
  is_x_max = x == max(x)  
)
```

```
#>   x y xy x2 y2 xy2 is_x_max  
#> 1 1 5  5  1 25   25  FALSE  
#> 2 2 6 12  4 36  144  FALSE  
#> 3 3 7 21  9 49  441  FALSE  
#> 4 4 8 32 16 64 1024  TRUE
```

Notice `mutate()` returns the original *and* new columns.

`mutate()` VS. `transmute()`

As their names imply, `mutate()` and `transmute()` are very similar functions.

- `mutate()` returns the **original** and **new** columns (variables).
- `transmute()` returns only the **new** columns (variables).

Note Both functions return a new object as *output*—they do not update the object in R's memory. (This is the case for all functions in `dplyr`.)

Pipes

We can't go much deeper into the land of `dplyr` without mentioning pipes.

A *pipe* in programming allows you to take the output of one function and plug it into another function as an argument/input.

In `dplyr`, the expression for a pipe is `%>%`.

R's pipe specifically plugs the returned object to the **left** of the pipe into the first argument of the function on the **right** fo the pipe, *e.g.*,

```
rnorm(10) %>% mean()
```

```
#> [1] 0.4854731
```

Pipes

Pipes help avoid lots of nested functions, prevent excessive writing to your disc, and increase the readability of our R scripts.

Example Three ways to draw 100 $N(0,1)$ observations and calculate the interquartile range (IQR: difference between the 75th and 25th percentiles).

```
# Save each intermediate step
draw <- rnorm(100)
end_points <- quantile(draw, probs = c(0.25, 0.75))
diff(end_points)

# Lots of nesting
diff(quantile(rnorm(100), probs = c(0.25, 0.75)))

# Piping 🍷
rnorm(100) %>% quantile(probs = c(0.25, 0.75)) %>% diff()
```

Pipes

By default, R pipes the output from the LHS of the pipe into the **first** argument of the function on the RHS of the pipe.

E.g., `a %>% fun(3)` is equivalent to `fun(arg1 = a, arg2 = 3)`.

If you want to pipe output into a different argument, you use a period (`.`).

- `b %>% fun(arg1 = 3, .)` is equivalent to `fun(arg1 = 3, arg2 = b)`.
- `b %>% fun(3, .)` is also equivalent to `fun(arg1 = 3, arg2 = b)`.
- `b %>% fun(., .)` is equivalent to `fun(arg1 = b, arg2 = b)`.

The `magrittr` package contains even more piping power.[†]

[†] `magrittr` = Magritte (of *this is not a pipe* fame) plus R.

dplyr

`%>%` and `dplyr`

Each `dplyr` function begins with a `.data` argument so that you can easily pipe in data frames (recall: `mutate(.data, ...)`).

The common workflow in `dplyr` will look something like

```
new_df ← old_df %>% mutate(cool stuff here)
```

which takes `old_df`, does some cool stuff with `mutate()`, and then saves the output of `mutate()` as `new_df`.

dplyr

`filter()`

The `filter()` function does what its name implies: it **filters the rows** of your data frame **based upon logical conditions**.

dplyr

filter()

The `filter()` function does what its name implies: it **filters the rows** of your data frame **based upon logical conditions**.

Example

```
# Create a dataset
some_df <- data.frame(
  x = 1:10,
  y = 11:20
)
```

```
# Only keep rows where x is 3
some_df %>% filter(x == 3)
```

```
#>   x  y
#> 1 3 13
```


dplyr

filter()

The `filter()` function does what its name implies: it **filters the rows** of your data frame **based upon logical conditions**.

Example

```
# Create a dataset
some_df <- data.frame(
  x = 1:10,
  y = 11:20
)
```

```
# Only keep rows where x > 7
some_df %>% filter(x > 7)
```

```
#>      x  y
#> 1   8 18
#> 2   9 19
#> 3  10 20
```

dplyr

`filter()`

The `filter()` function does what its name implies: it **filters the rows** of your data frame **based upon logical conditions**.

Example

```
# Create a dataset
some_df <- data.frame(
  x = 1:10,
  y = 11:20
)
```

```
# Keep rows where y/x > 3
some_df %>% filter(y/x > 3)
```

```
#>   x  y
#> 1  1 11
#> 2  2 12
#> 3  3 13
#> 4  4 14
```

dplyr

`filter()`

The `filter()` function does what its name implies: it **filters the rows** of your data frame **based upon logical conditions**.

Example

```
# Create a dataset
some_df <- data.frame(
  x = 1:10,
  y = 11:20
)
```

```
# Keep rows where x>7 OR y<12
some_df %>%
  filter(x > 7 | y < 12)
```

```
#>      x  y
#> 1   1 11
#> 2   8 18
#> 3   9 19
#> 4  10 20
```

dplyr

filter()

The `filter()` function does what its name implies: it **filters the rows** of your data frame **based upon logical conditions**.

Example

```
# Create a dataset
some_df <- data.frame(
  x = 1:10,
  y = 11:20
)
```

```
# Keep rows where  $15 \leq y \leq 18$ 
some_df %>%
  filter(between(y, 15, 18))
```

```
#>   x  y
#> 1  5 15
#> 2  6 16
#> 3  7 17
#> 4  8 18
```

dplyr

`filter()`

The `filter()` function does what its name implies: it **filters the rows** of your data frame **based upon logical conditions**.

Example

```
# Create a dataset
some_df <- data.frame(
  x = 1:10,
  y = 11:20
)
```

```
# Keep rows where y > 20
some_df %>% filter(y > 20)
```

```
#> [1] x y
#> <0 rows> (or 0-length row.names)
```

If you filter your data frame down to nothing, R returns a 0-row data frame with the names/number of columns from the original data frame.

dplyr

`select()`

Just as `filter()` grabs row-based subsets of your data frame, `select()` grabs column-based subsets.

You can select columns using their **names**

```
our_df %>% select(var10, var100)
```

you can select columns using their **numbers**

```
our_df %>% select(10, 100)
```

or you can select columns using **helper functions**

```
our_df %>% select(starts_with("var10"))
```

`select()` helps you narrow down a dataset to its necessary features.

`summarize()`

Hopefully you're starting to see that functions' names in `dplyr` tell you what the function does.

`summarize()`[†] summarizes variables—you choose the variables and the summaries (e.g., `mean()` or `min()`).

```
the_df %>% summarize(  
  mean(x), mean(y), mean(z),  
  min(x), max(x),  
)
```

would return a 1×5 data frame with the means of `x`, `y`, and `z`; the minimum of `x`; and the maximum of `x`.

[†] or `summarise()` if you ❤️ 🇬🇧

dplyr

`summarize()` and `group_by()`

While sample-wide summaries are certainly interesting, `dplyr` has one last gem for us: `group_by()`.

`group_by()` groups your observations by the variable(s) that you name.

Specifically, `group_by()` returns a *grouped data frame* that you can then feed to `summarize()`, `mutate()`, or `transmute` to perform grouped calculations, e.g., each group's mean.

Example: Grouped summaries

```
# Create a new data frame
our_df <- data.frame(
  x = 1:6,
  y = c(0, 1),
  grp = rep(c("A", "B"), each = 3)
)
```

```
#>   x y grp
#> 1 1 0  A
#> 2 2 1  A
#> 3 3 0  A
#> 4 4 1  B
#> 5 5 0  B
#> 6 6 1  B
```

```
# For dataset 'our_df' ...
our_df %>%
  # Group by 'grp'
  group_by(grp) %>%
  # Take means of 'x' and 'y'
  summarize(mean(x), mean(y))
```

```
#> # A tibble: 2 x 3
#>   grp   `mean(x)` `mean(y)`
#>   <fct>     <dbl>     <dbl>
#> 1 A             2     0.333
#> 2 B             5     0.667
```

Example: Grouped mutation

```
# Create a new data frame
our_df <- data.frame(
  x = 1:6,
  y = c(0, 1),
  grp = rep(c("A", "B"), each = 3)
)
```

```
#>   x y grp
#> 1 1 0  A
#> 2 2 1  A
#> 3 3 0  A
#> 4 4 1  B
#> 5 5 0  B
#> 6 6 1  B
```

```
# Add grp means for x and y
our_df %>%
  group_by(grp) %>%
  mutate(
    x_m = mean(x), y_m = mean(y)
  )
```

```
#> # A tibble: 6 x 5
#> # Groups:   grp [2]
#>       x     y grp     x_m     y_m
#>   <int> <dbl> <fct> <dbl> <dbl>
#> 1     1     0  A         2  0.333
#> 2     2     1  A         2  0.333
#> 3     3     0  A         2  0.333
#> 4     4     1  B         5  0.667
#> 5     5     0  B         5  0.667
#> 6     6     1  B         5  0.667
```

dplyr

arrange()

`arrange()` will sort the rows of a data frame using the inputted columns.

R defaults to starting with the "lowest" (smallest) at the top of the data frame. Use a `-` in front of the variable's name to reverse sort.

```
# As is  
our_df
```

```
#>   x y grp  
#> 1 1 0   A  
#> 2 2 1   A  
#> 3 3 0   A  
#> 4 4 1   B  
#> 5 5 0   B  
#> 6 6 1   B
```

```
# As is  
our_df %>% arrange(y, grp, -x)
```

```
#>   x y grp  
#> 1 3 0   A  
#> 2 1 0   A  
#> 3 5 0   B  
#> 4 2 1   A  
#> 5 6 1   B  
#> 6 4 1   B
```

Table of contents

Data and R

1. Schedule
2. Matrix review
3. The `data.frame`
 - Basic examples
 - Creating
 - Indexing
 - Names
 - Adding variables

dplyr

1. Intro
2. `mutate()`
3. `transmute()`
4. Pipes (`%>%`)
5. `filter()`
6. `select()`
7. `summarize`
8. `summarize()` and `group_by()`
9. `arrange()`