

Plotting in R

EC 425/525, Lab 5

Edward Rubin

09 May 2019

Prologue

Schedule

Last time

Regression

Today

Plotting in R (especially `ggplot2`)

Plotting

Plotting

The default option: `plot()`

While we'll quickly move on to other options, R's `plot()` function (in the default `graphics` package) is a great tool for basic data exploration—it's fast, simple, and flexible.

In fact, `plot()` is a generic function, that works for many classes.

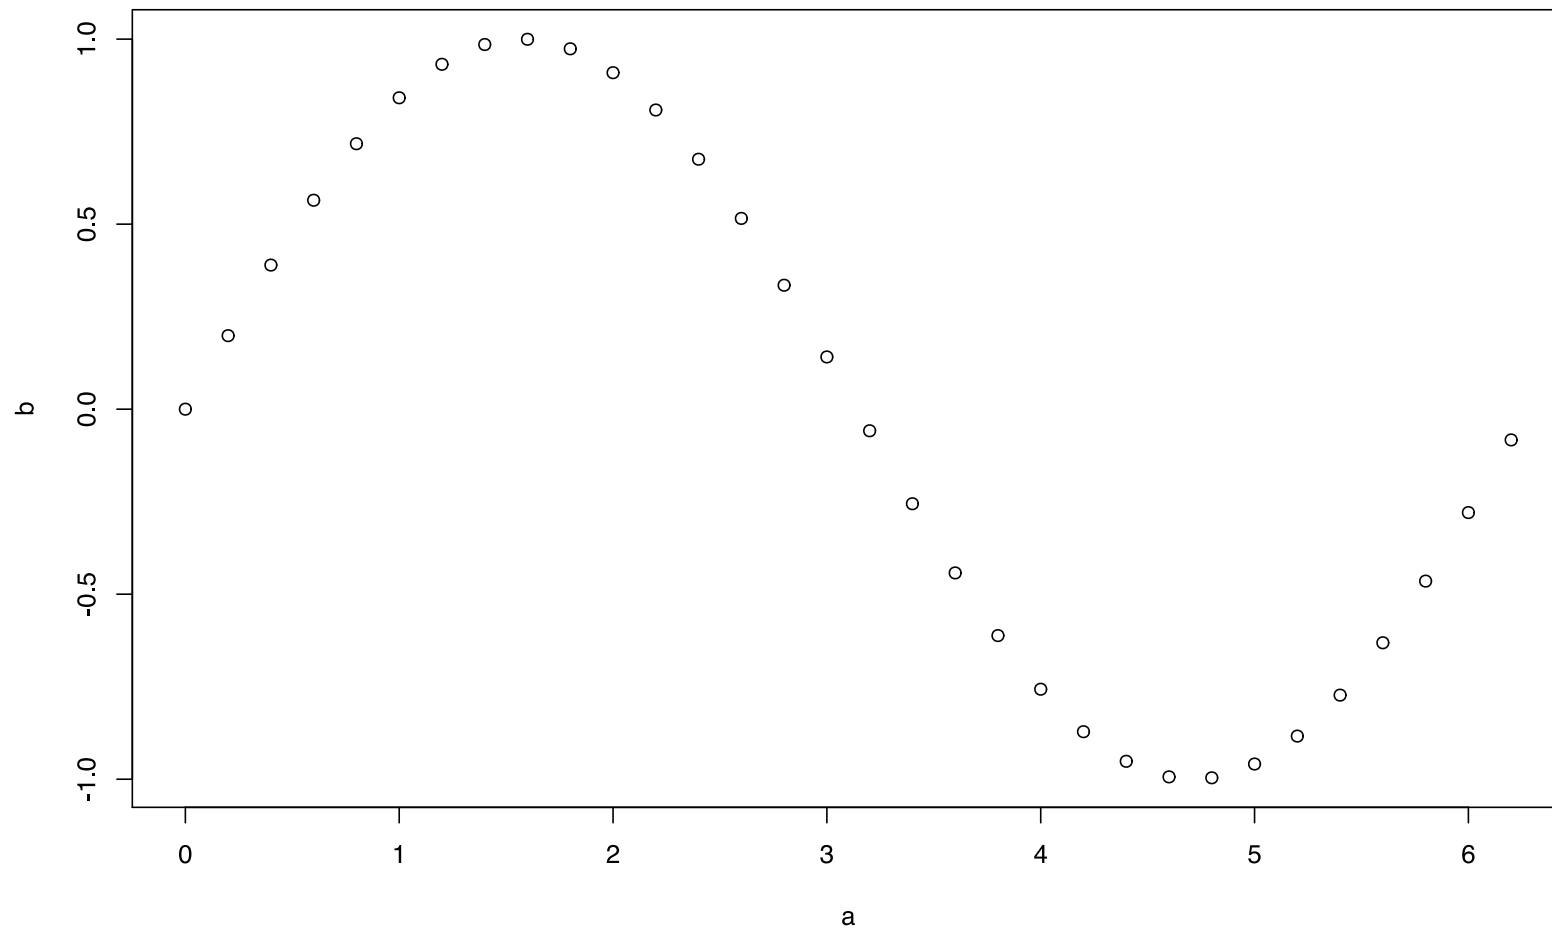
General arguments for `plot()`:

- `x` and `y` for coordinates
- `type = { "p" oints, "l" ines, etc.}` (*optional*)
- `xlab`, `ylab`, `main`, and `sub` for axis labels and (sub)title (*optional*)
- `col` and `pch` for color and plot character (*optional*)
- `lty` and `lwd` for line type, and line width (*optional*)

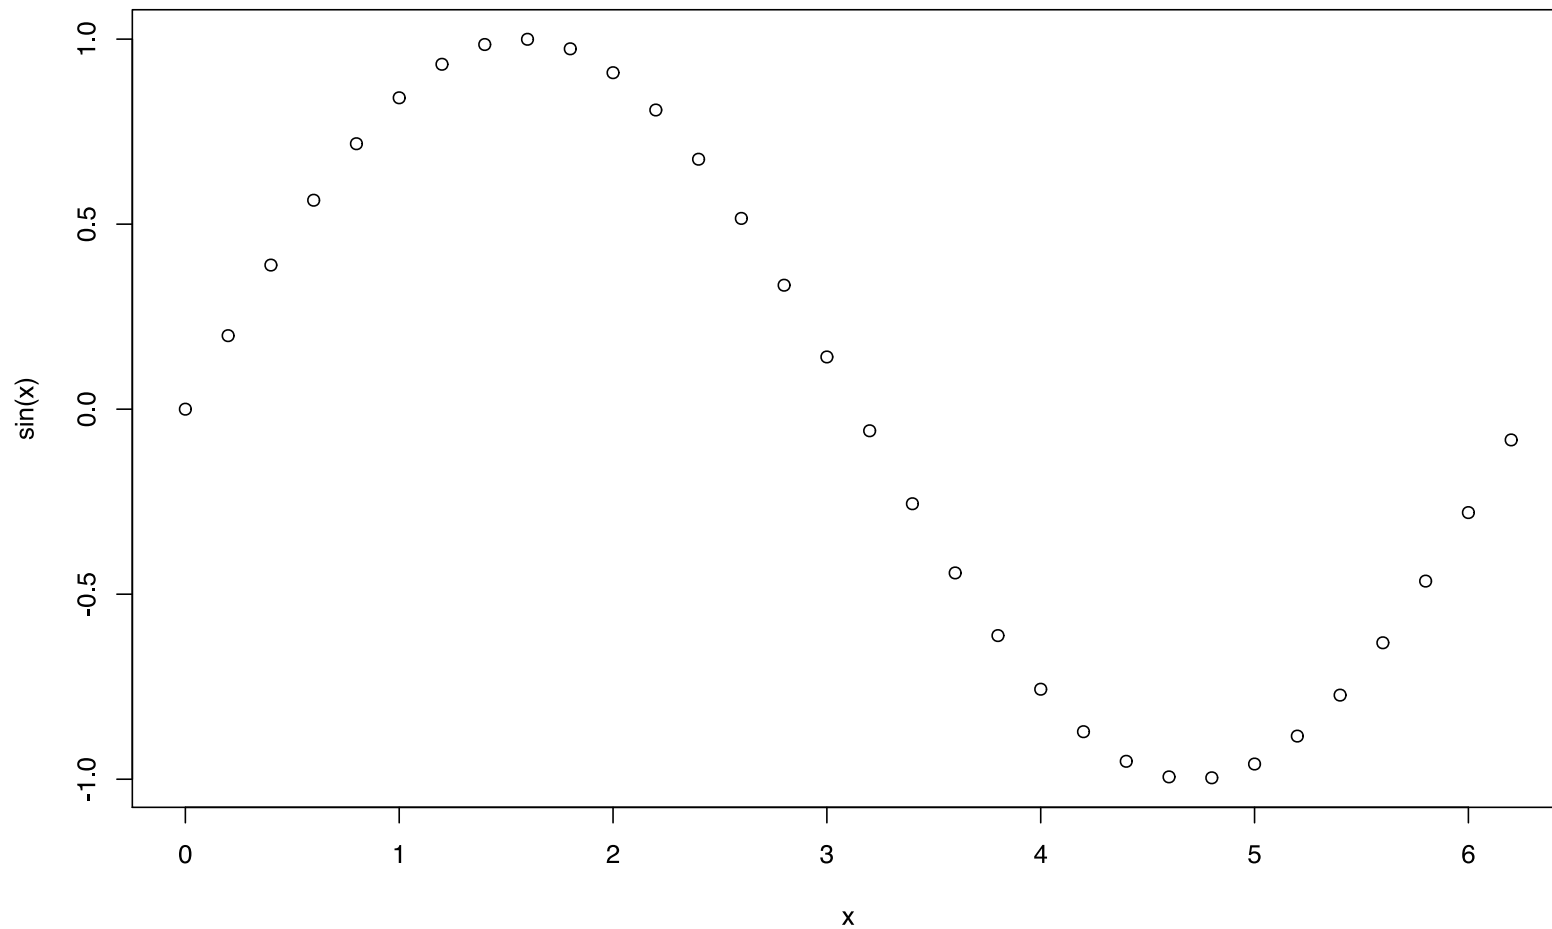
Let's see `plot()` in action.

```
# Define two vectors  
a ← seq(from = 0, to = 2*pi, by = 0.2)  
b ← sin(a)
```

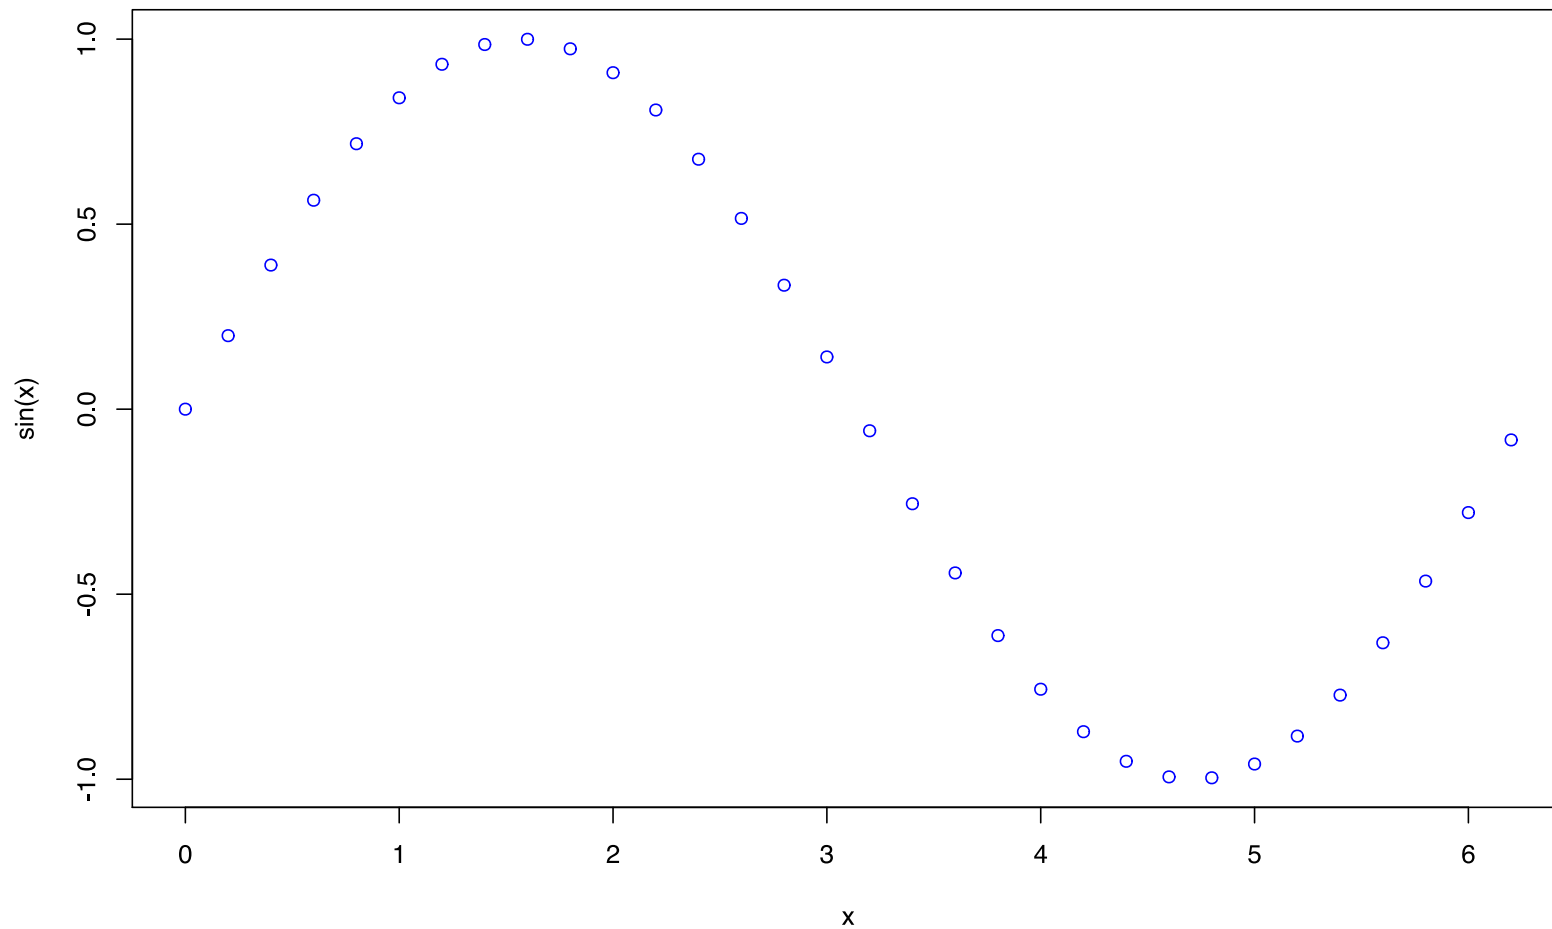
```
plot(x = a, y = b)
```



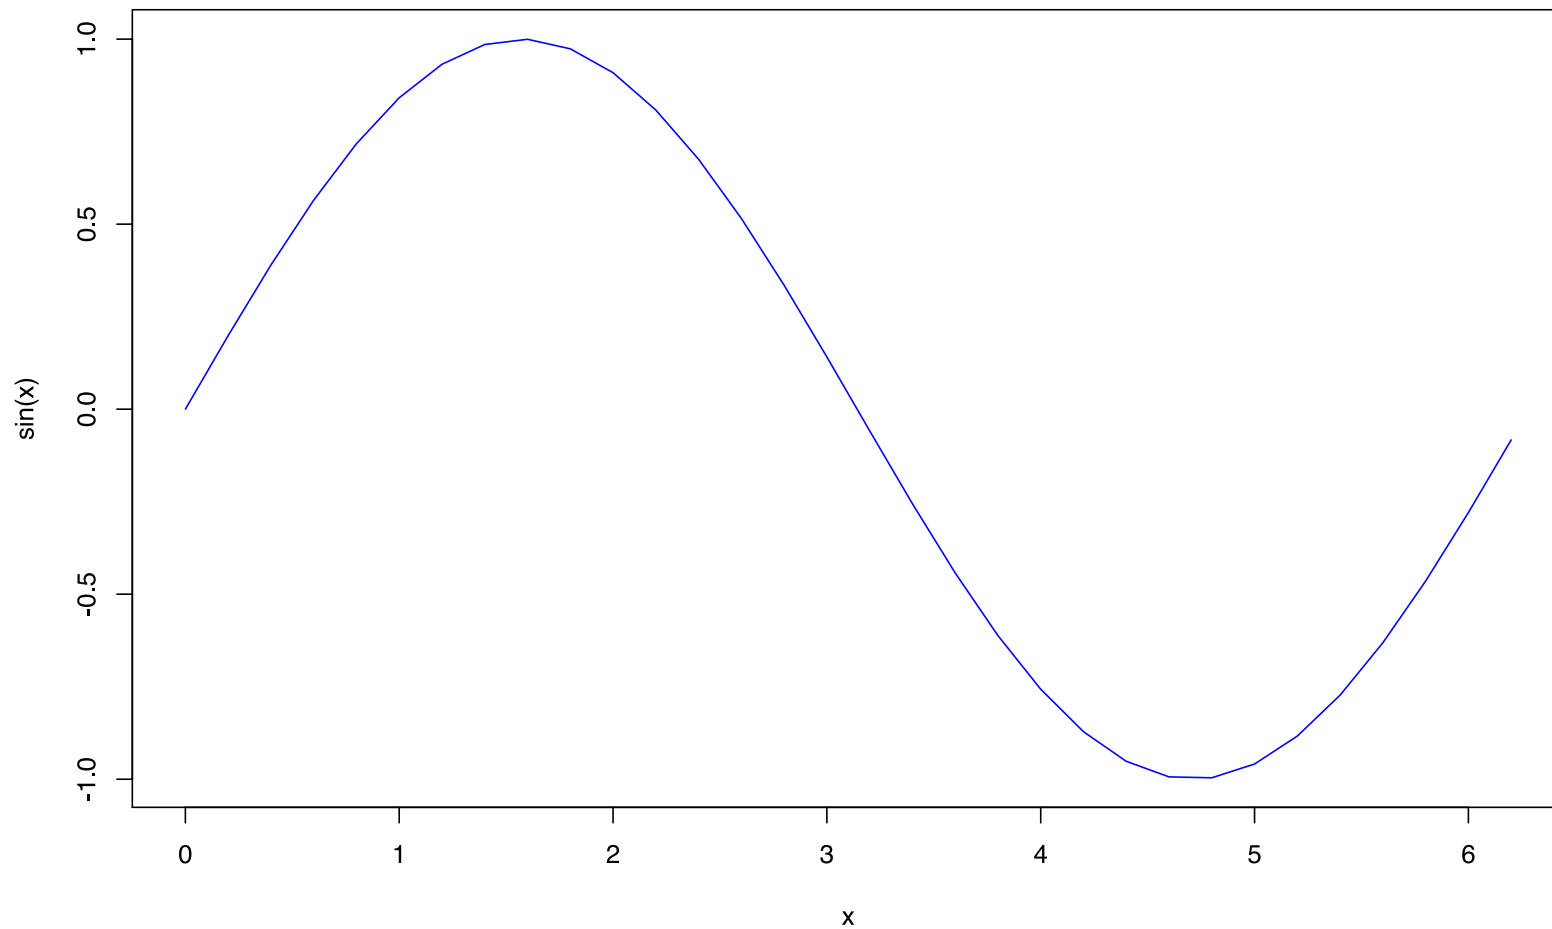
```
plot(x = a, y = b, xlab = "x", ylab = "sin(x)")
```



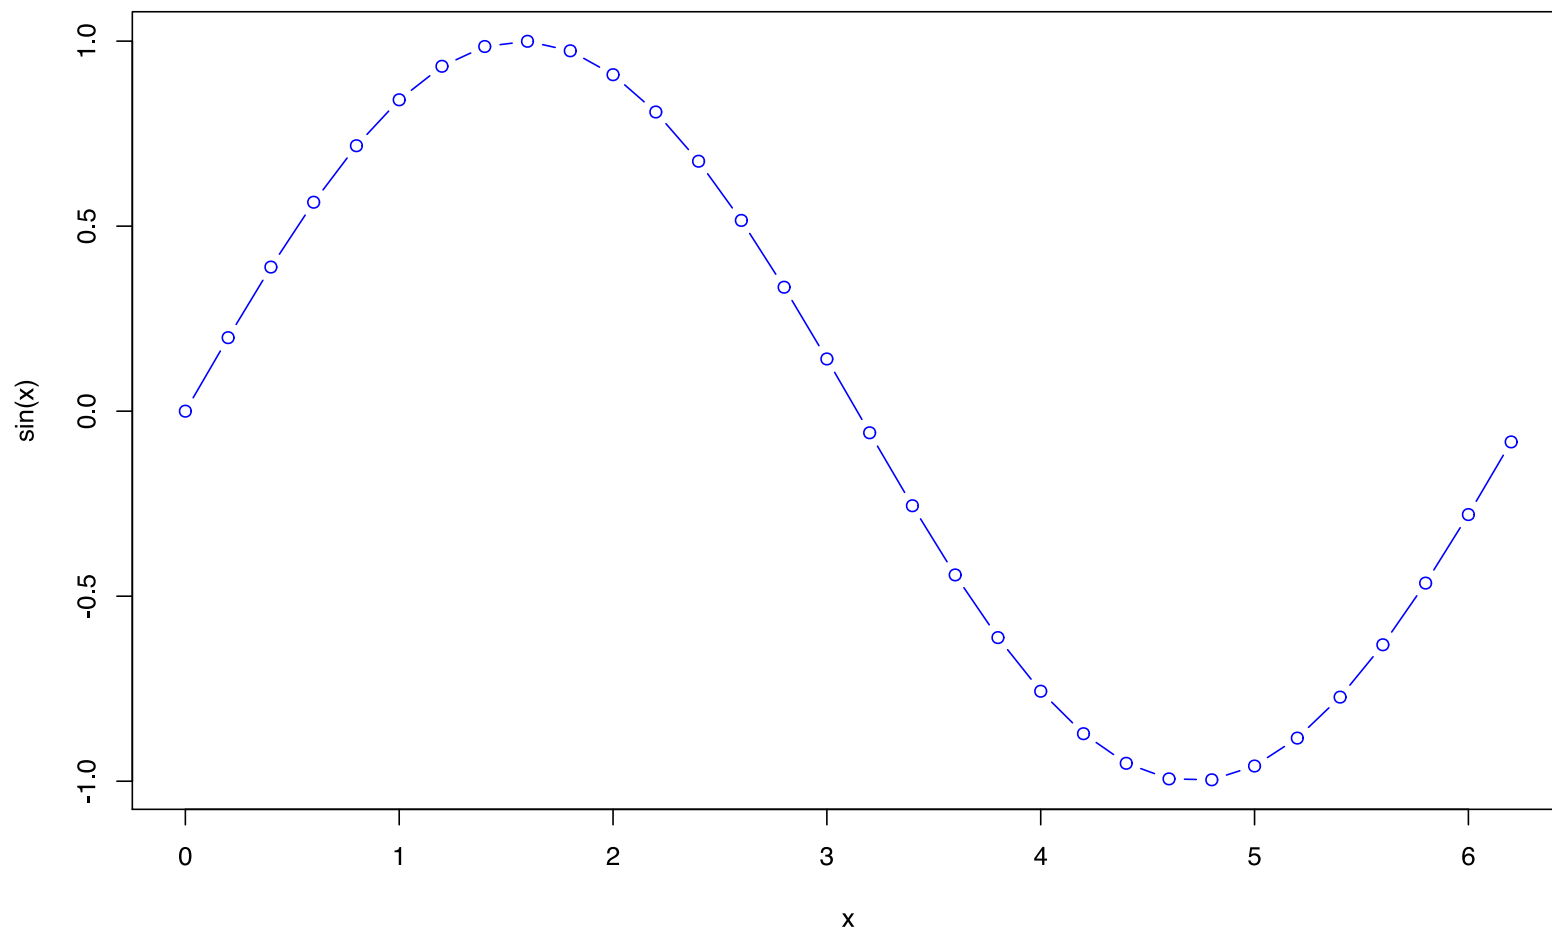

```
plot(x = a, y = b, xlab = "x", ylab = "sin(x)", col = "blue")
```



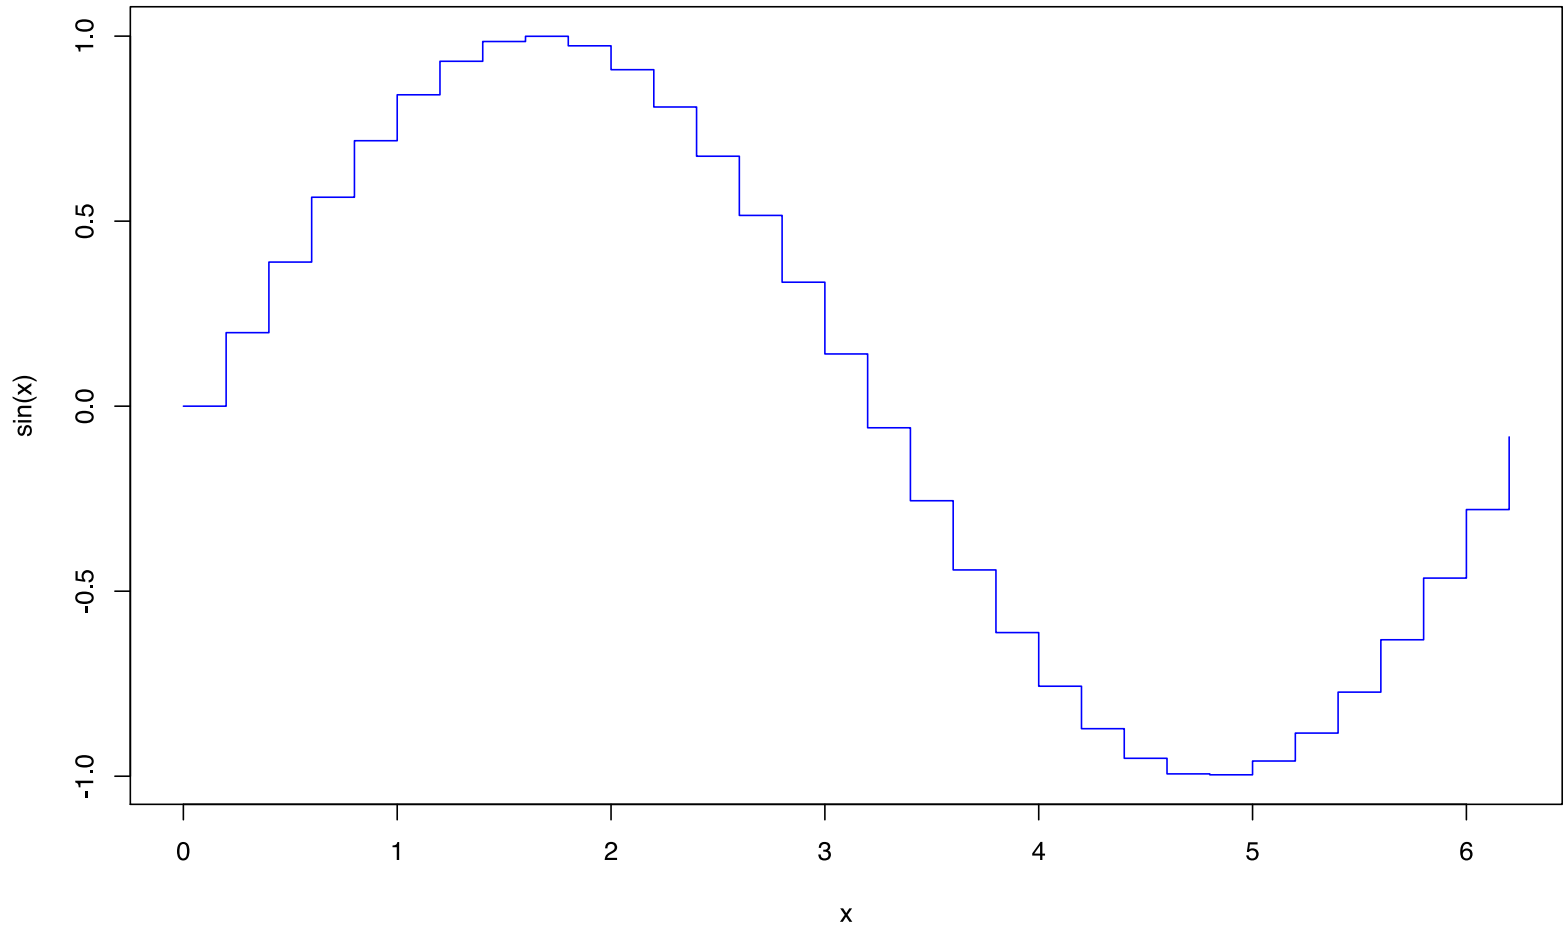
```
plot(x = a, y = b, xlab = "x", ylab = "sin(x)", col = "blue", type = "l")
```



```
plot(x = a, y = b, xlab = "x", ylab = "sin(x)", col = "blue", type = "b")
```



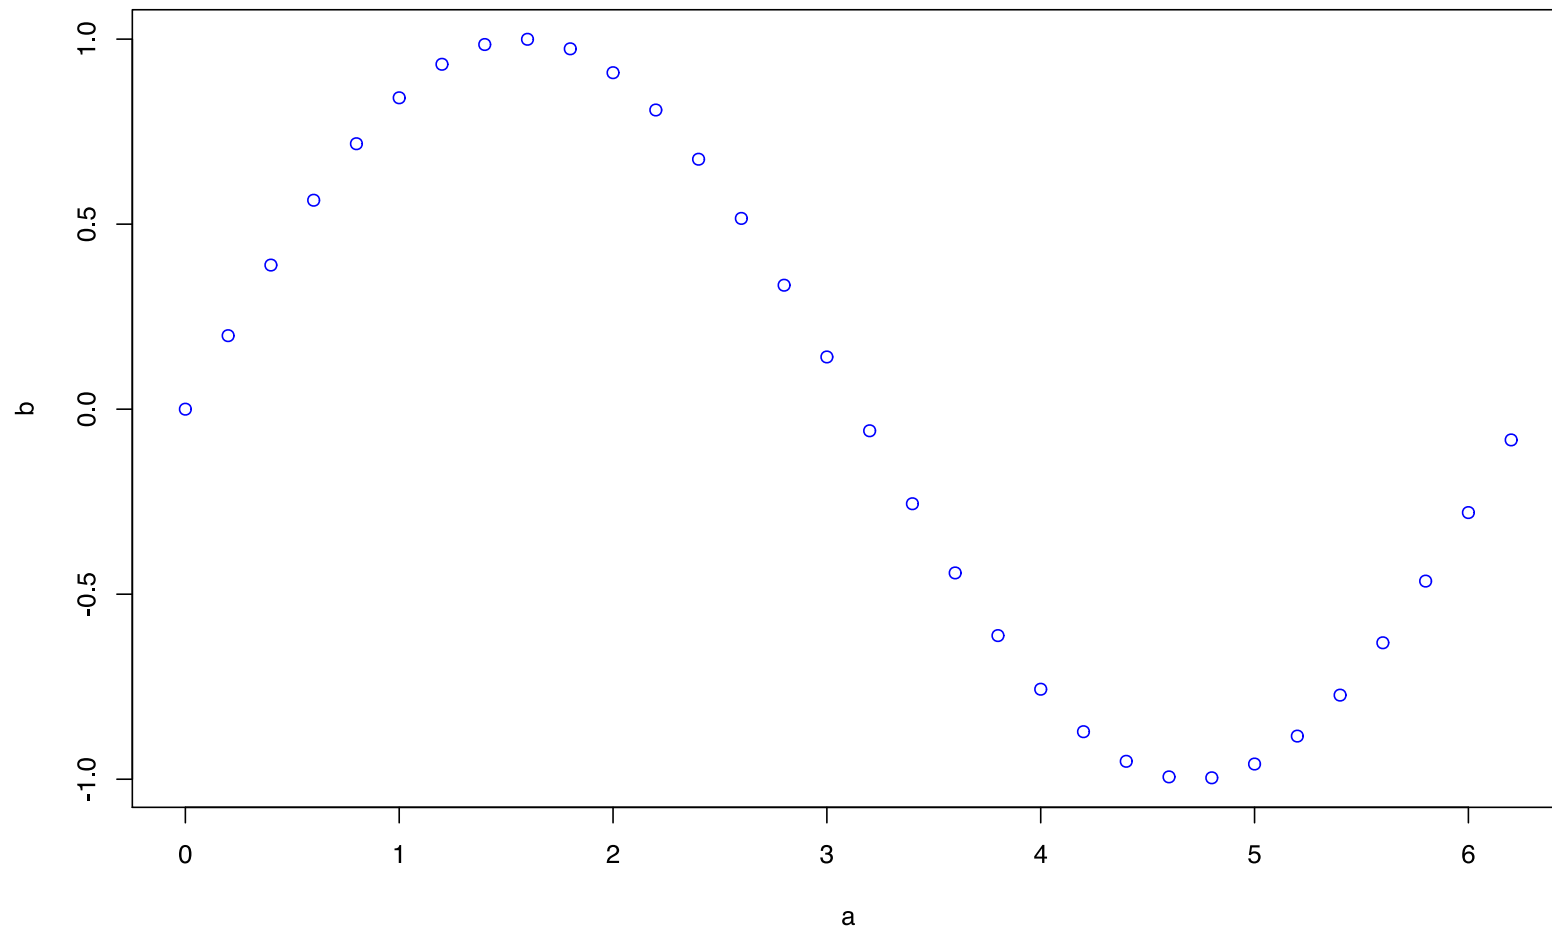
```
plot(x = a, y = b, xlab = "x", ylab = "sin(x)", col = "blue", type = "s")
```



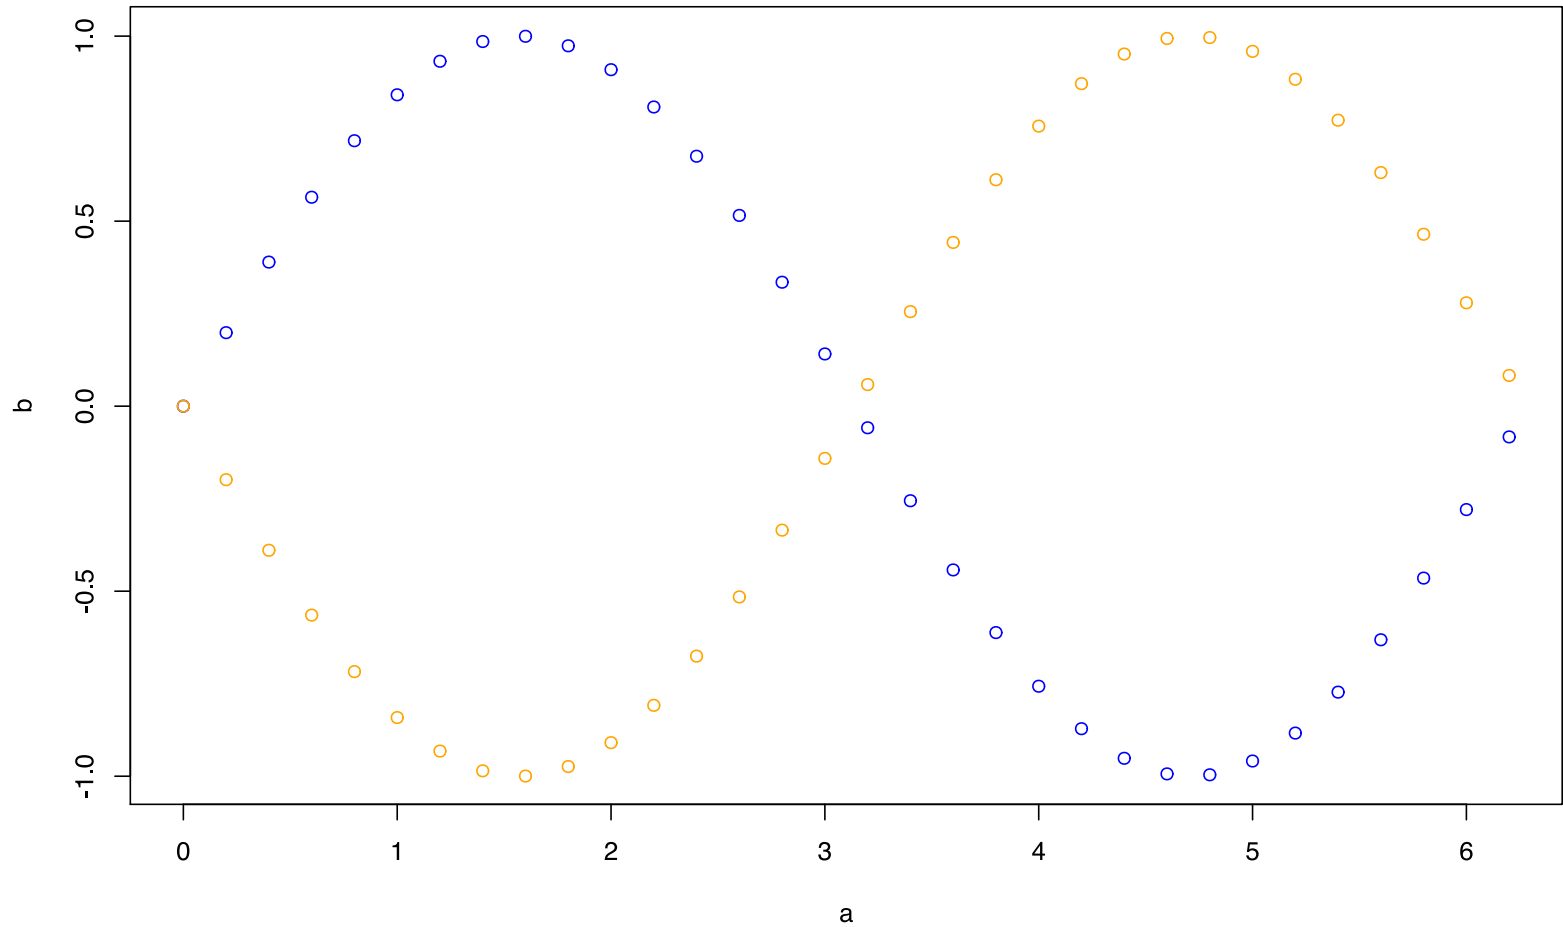
`plot()` is essentially calling `points()` or `lines()`.

You can layer plots by using these individual functions.

```
plot(x = a, y = b, col = "blue")
```

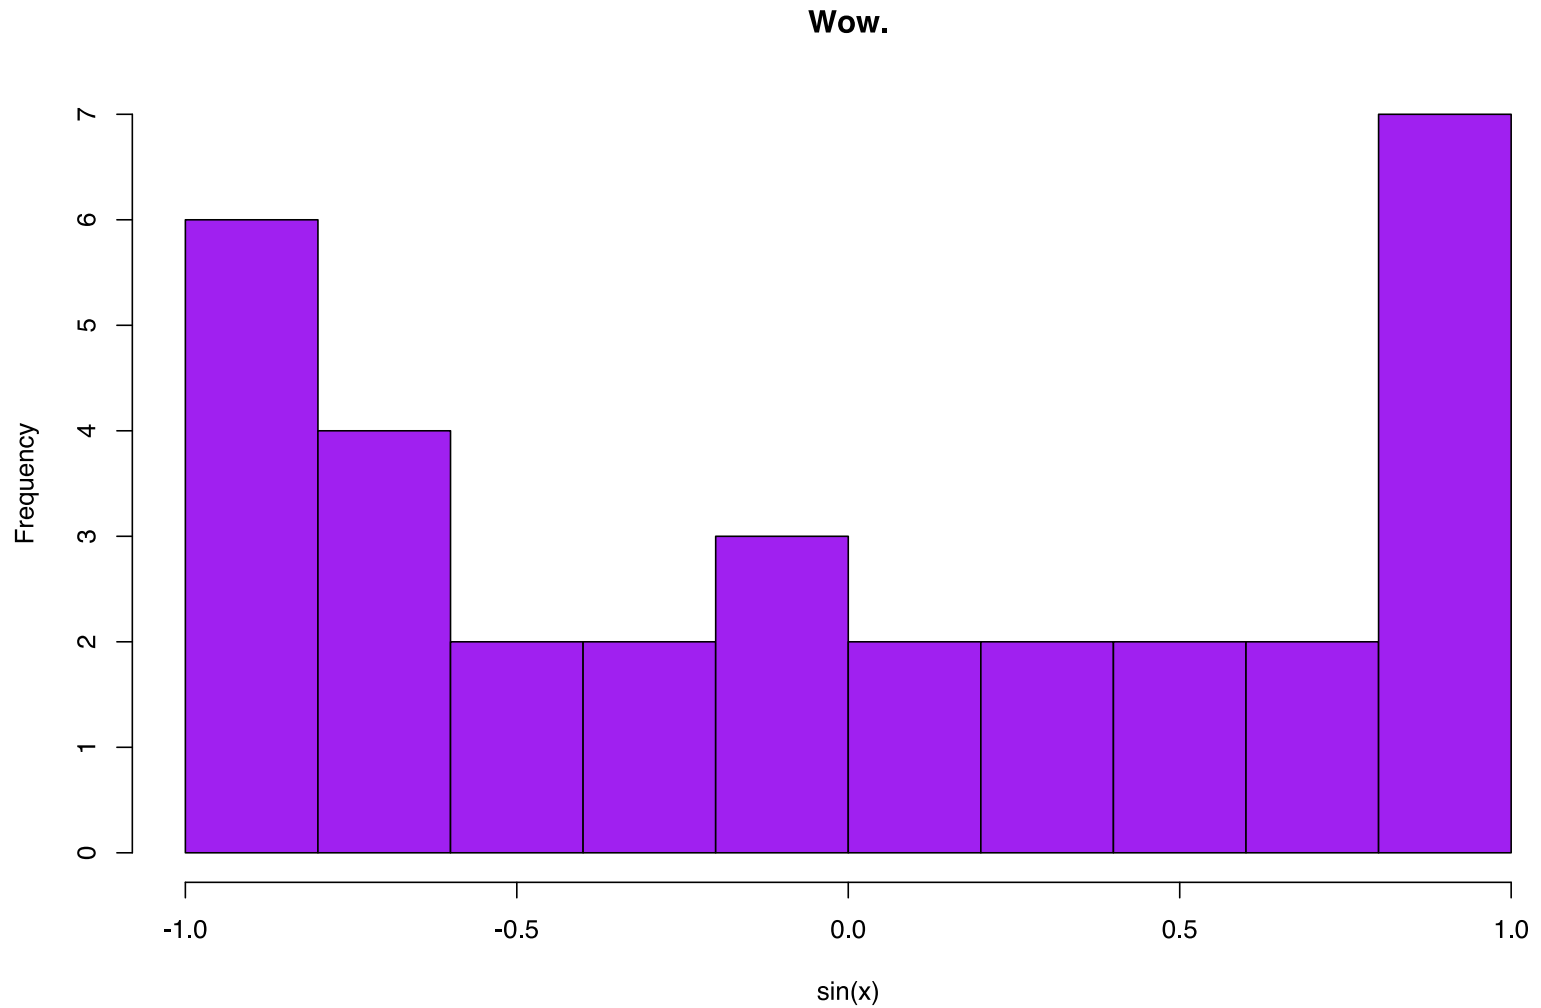


```
plot(x = a, y = b, col = "blue"); points(x = a, y = -b, col = "orange")
```



`graphics` also offers a nice histogram function in `hist()`.


```
hist(x = b, breaks = 10, col = "purple", xlab = "sin(x)", main = "Wow.")
```



That said/done, further customization/manipulation of your graphics using `graphics` plotting functions can become quite difficult.

Enter `ggplot2`

ggplot2

ggplot2

The grammar

The `ggplot2` package offers an incredibly flexible, diverse, and powerful set of functions for creating graphics in R.

The `gg` stands for the *grammar of graphics*.

`ggplot2`

1. centers on a **data frame** (the `data` argument)
2. maps variables to **aesthetics** (the `aes` argument)
3. **layers geometries** to *build up* your graphic

Note The package is called `ggplot2`, but the main function is `ggplot()`.

ggplot2

ggplot()

Main arguments

1. **data** Your dataset. As a data frame (or `tibble`).
1. **aes()** Maps variables in `data` to "aesthetics" like `x`, `color`, `shape`.

Example A time series of problems, `color` defined by money

```
library(ggplot2)
ggplot(
  data = pretend_df,
  aes(x = time, y = problems, color = money)
)
```

ggplot2

Layers

The `ggplot()` function doesn't plot anything—it *sets up* the plot.

To create the actual figure, you layer **geometries** (e.g., `geom_point()`), **scales** (e.g., `scale_color_manual()`), and other **options** (e.g., `xlab()`).

You **add layers** using the addition sign (+).

Example A time series of problems, `color` defined by money

```
library(ggplot2)
ggplot(
  data = pretend_df,
  aes(x = time, y = problems, color = money)
) +
  geom_point() + geom_line()
```

Alright, let's build a plot.

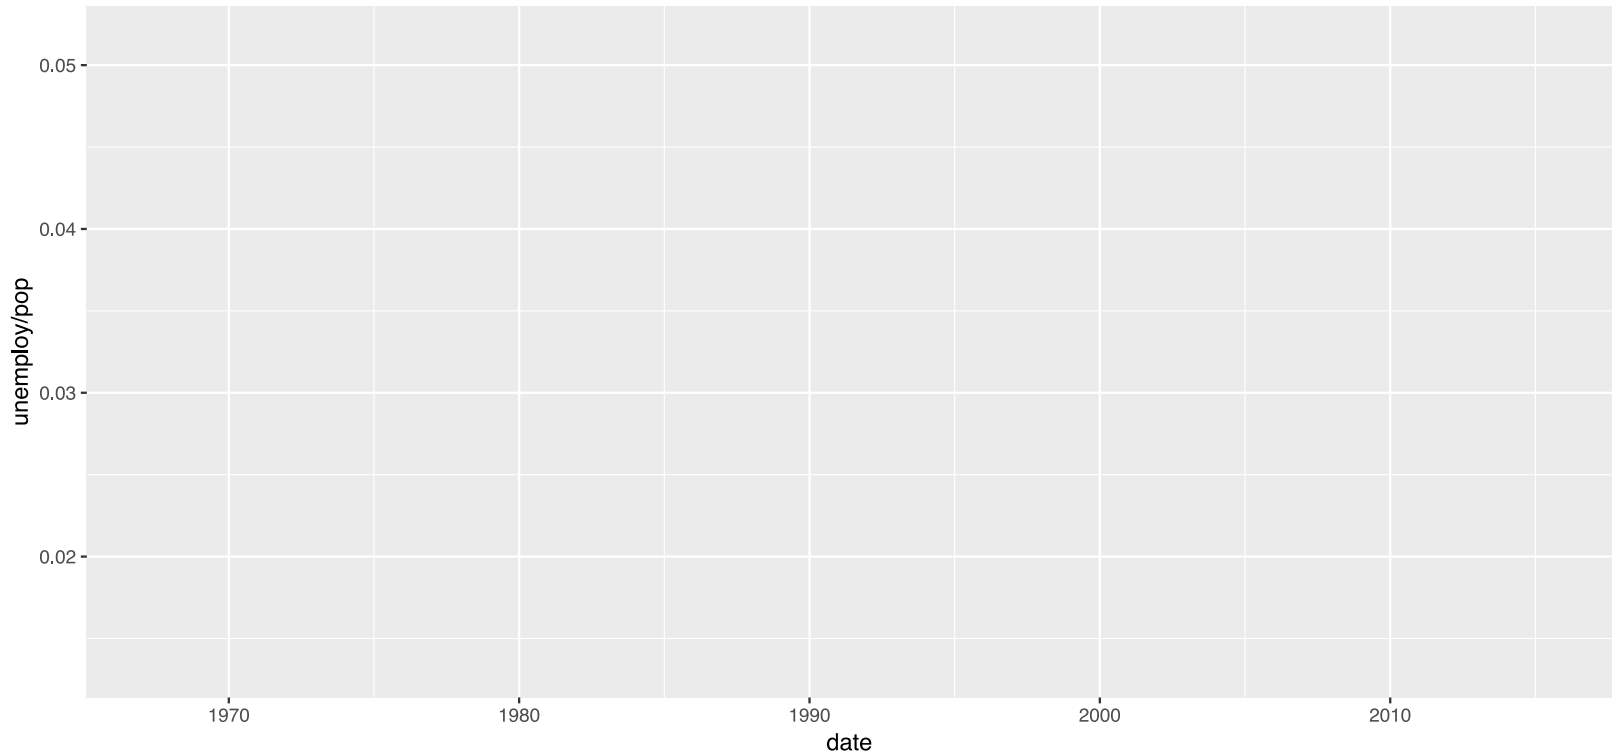
We'll use the `economics` dataset that comes with `ggplot2` (because economics).

	date	pce	pop	psavert	uempmed	unemploy
1	1967-07-01	507.4	198712	12.5	4.5	2944
2	1967-08-01	510.5	198911	12.5	4.7	2945
3	1967-09-01	516.3	199113	11.7	4.6	2958
4	1967-10-01	512.9	199311	12.5	4.9	3143
5	1967-11-01	518.1	199498	12.5	4.7	3066
6	1967-12-01	525.8	199657	12.1	4.8	3018
7	1968-01-01	531.5	199808	11.7	5.1	2878
8	1968-02-01	534.2	199920	12.2	4.5	3001

Showing 1 to 8 of 574 entries

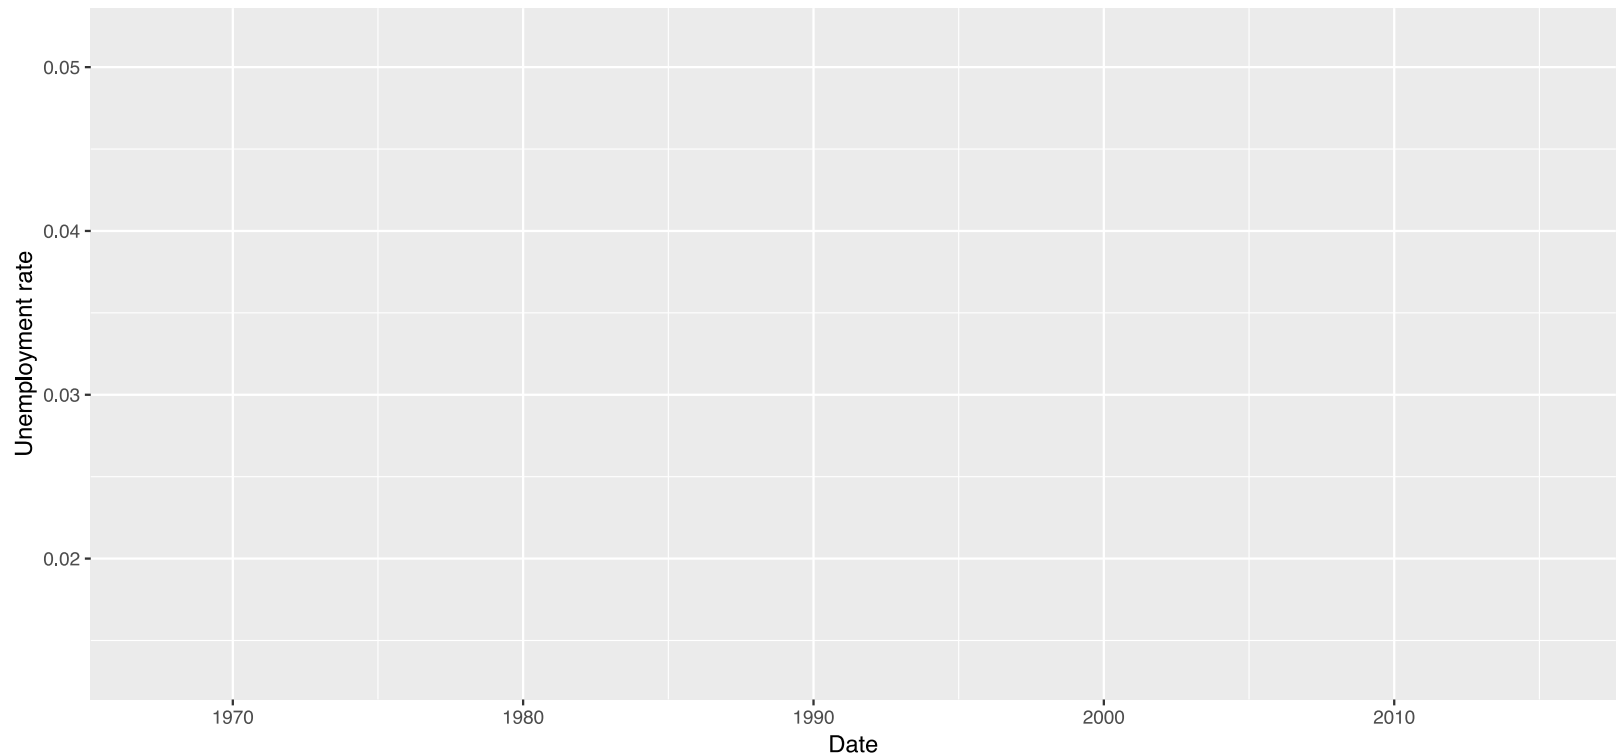
Set up the plot.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop))
```



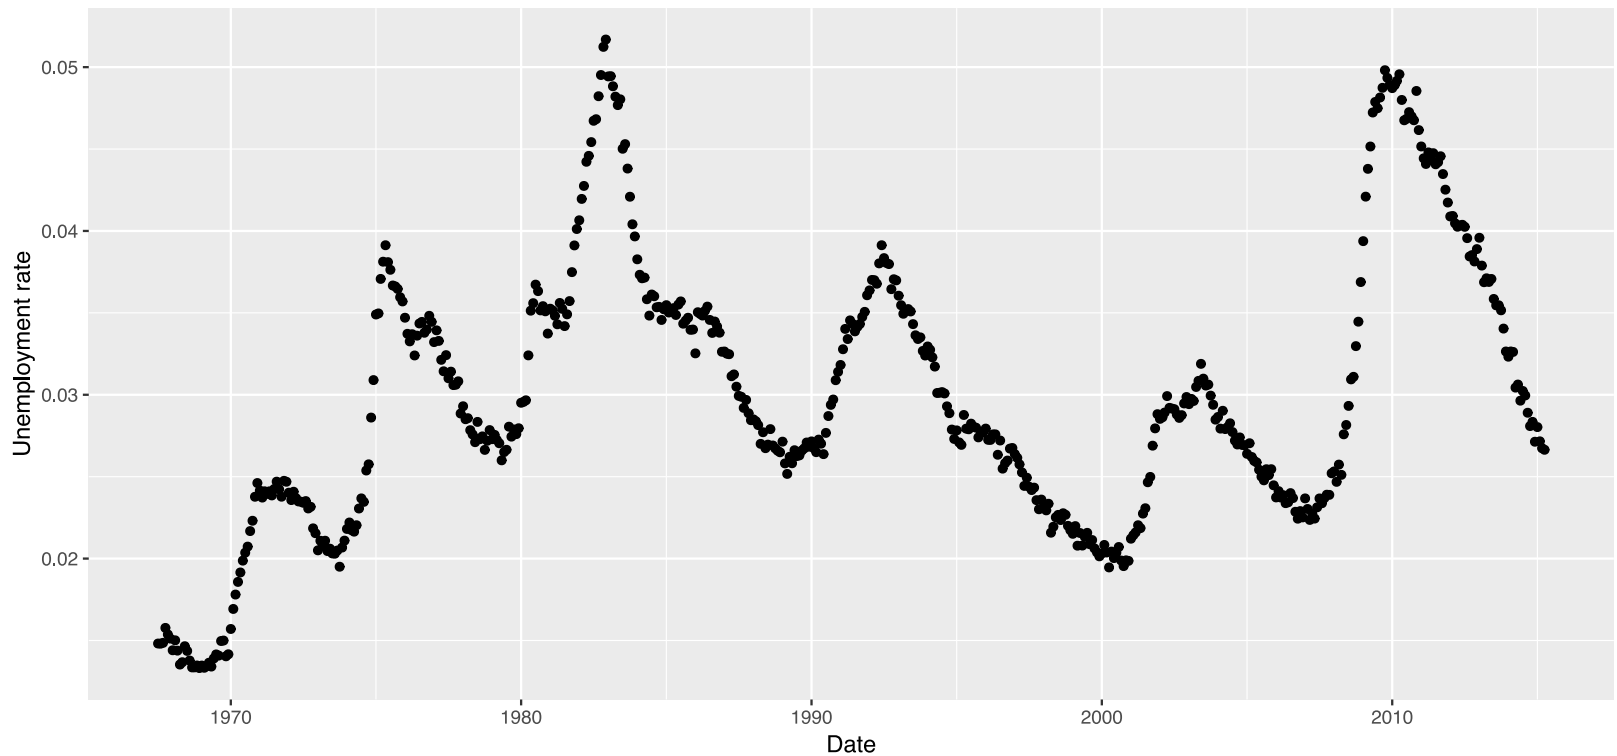
Label the axes.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date")
```



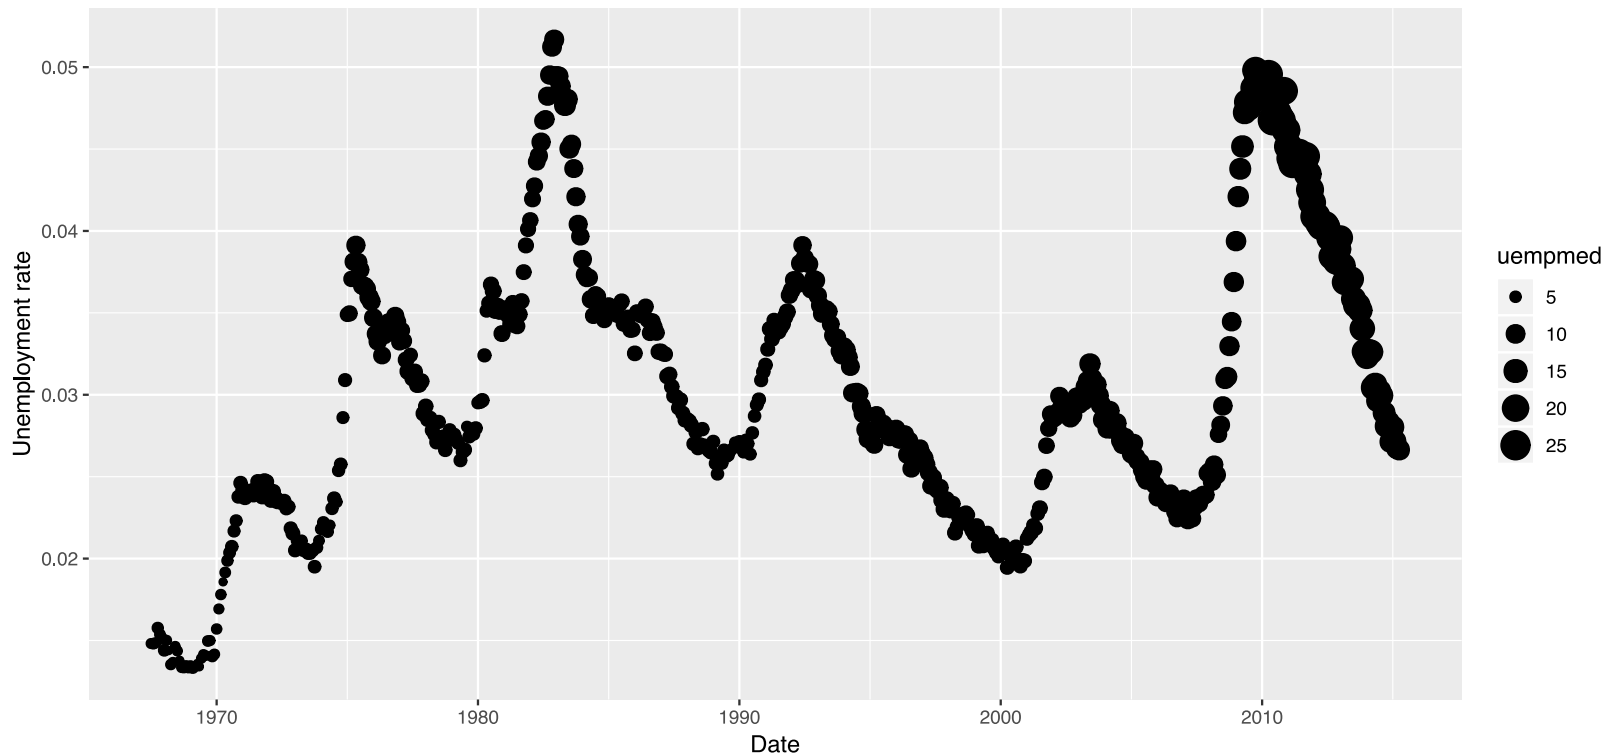
Draw some points.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point()
```



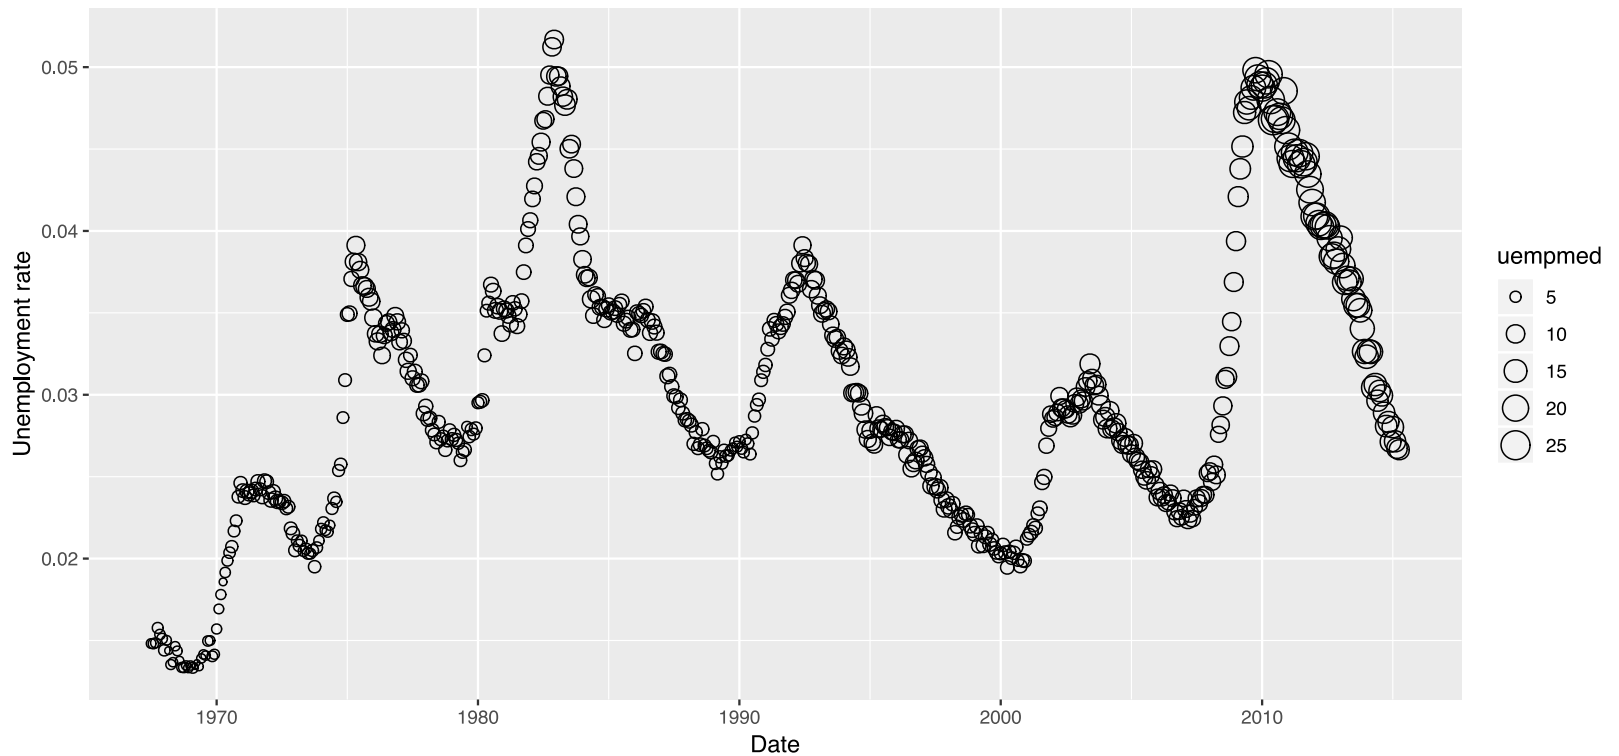
Map the `size` to the median duration of unemployment.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop, size = uempmed)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point()
```



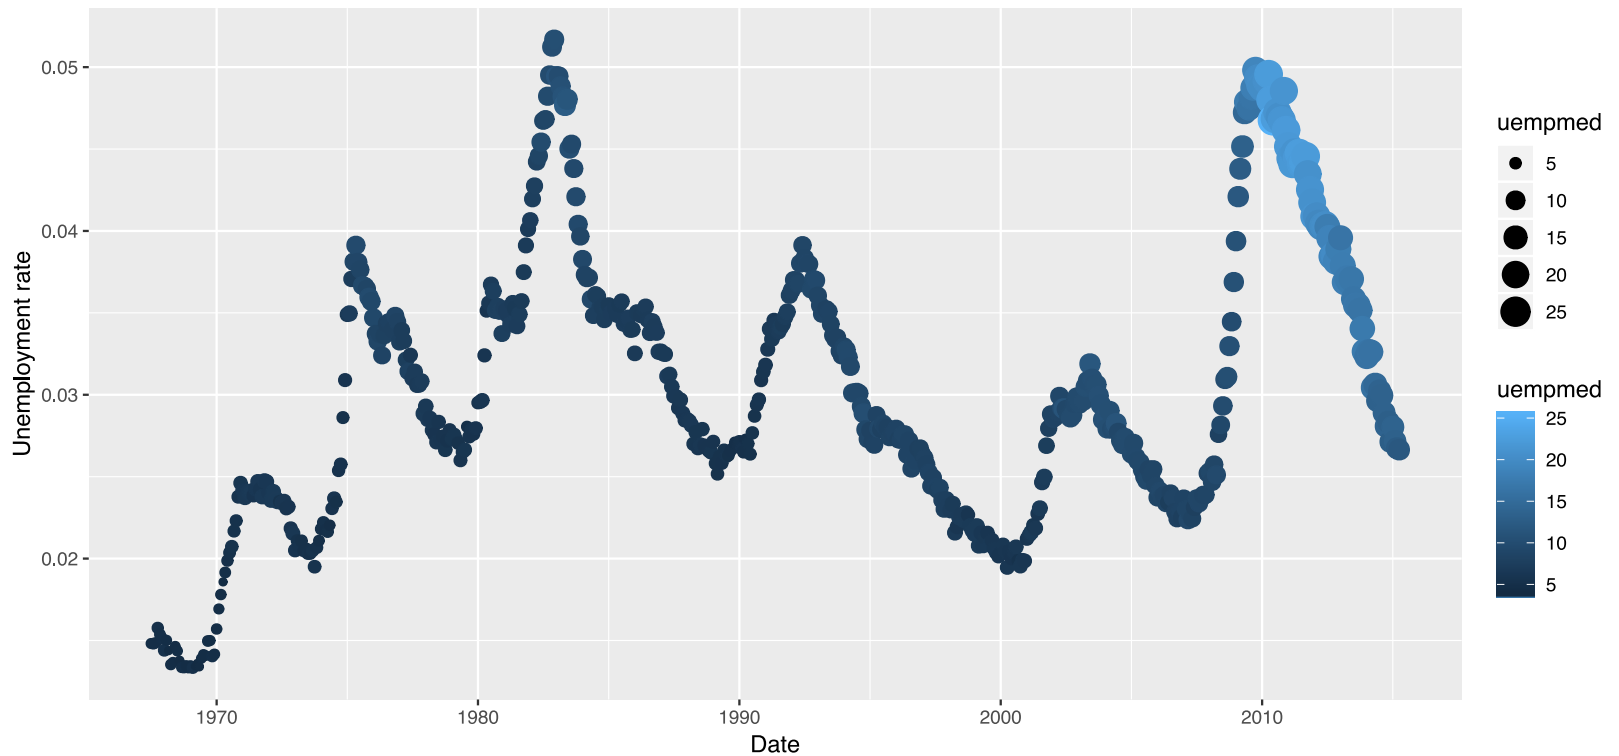
Change the `shape` of the points.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop, size = uempmed)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(shape = 1)
```



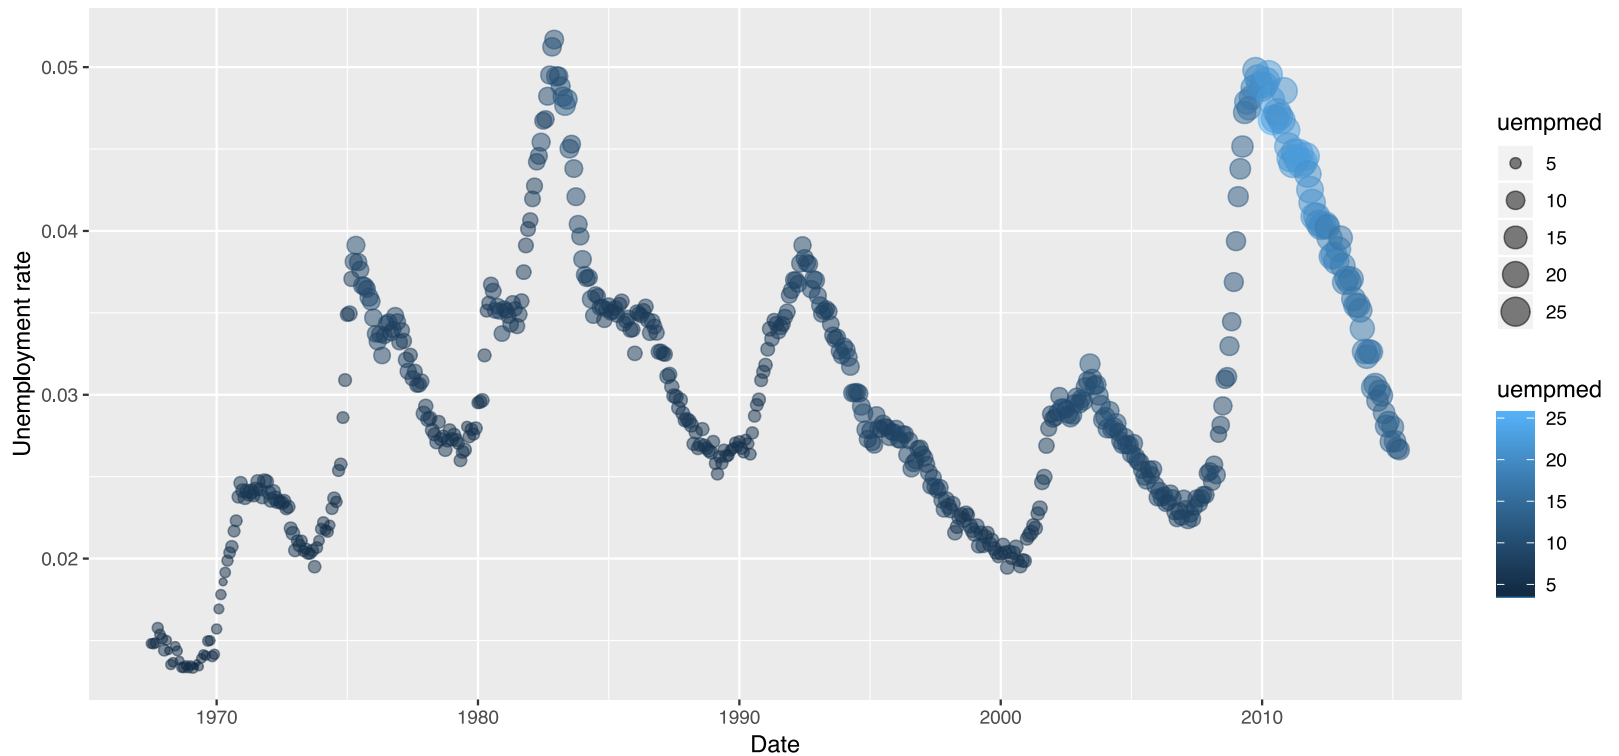
Map points' `color` to the median duration of unemployment.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop, size = uempmed)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed))
```



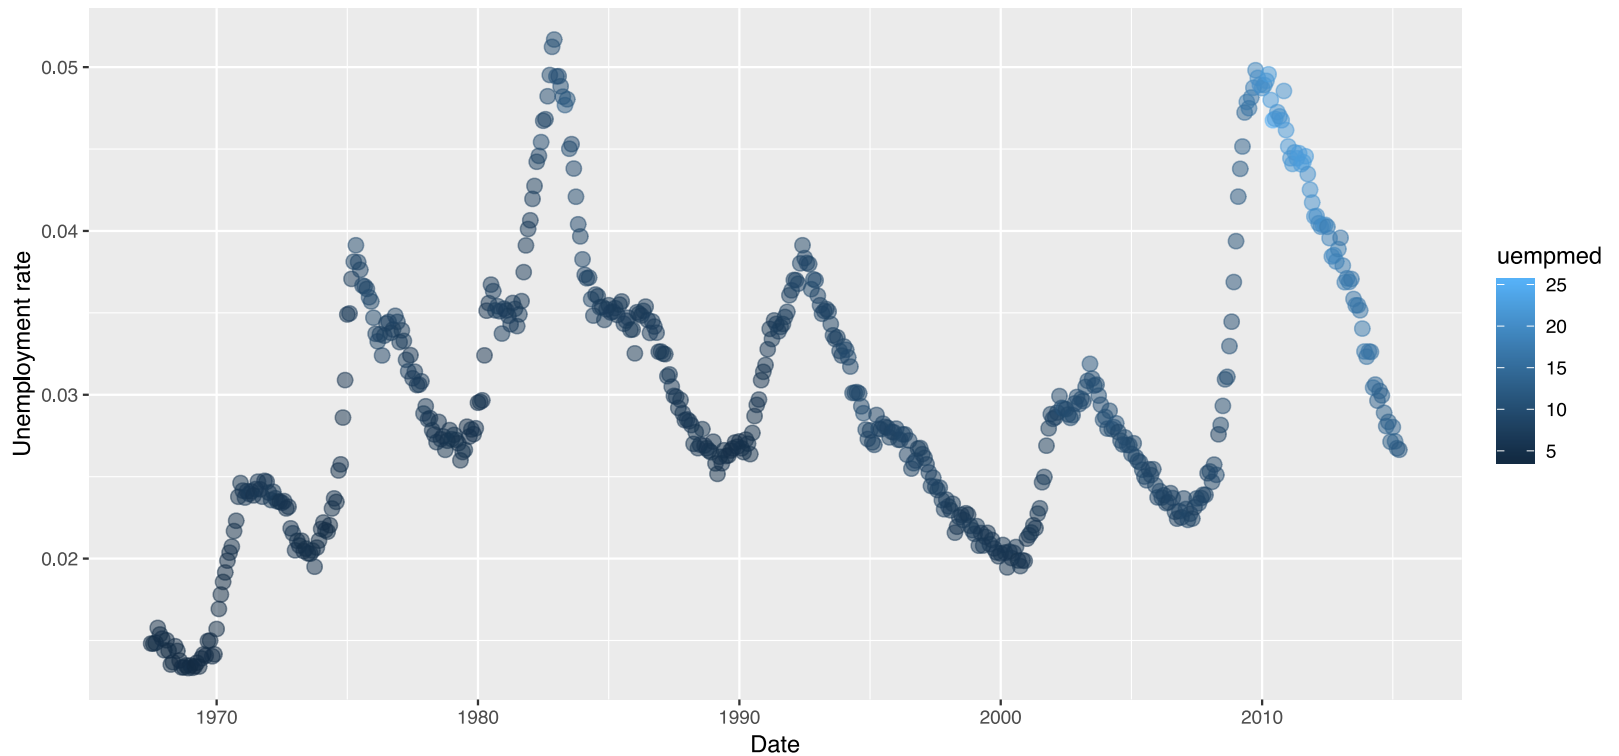
Add some transparency (`alpha`) to our points.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop, size = uempmed)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed), alpha = 0.5)
```



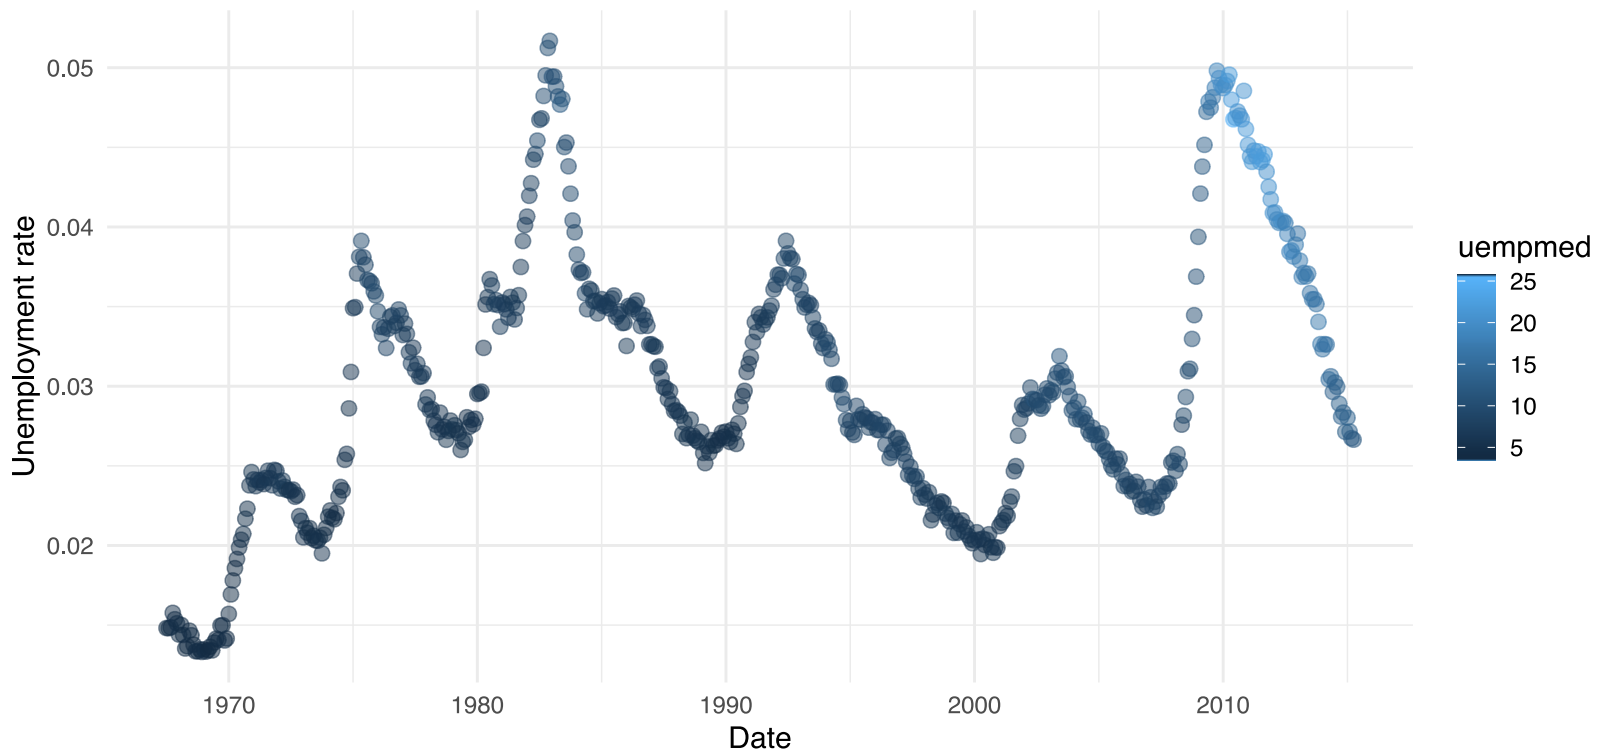
Same size points; all bigger.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed), alpha = 0.5, size = 3)
```



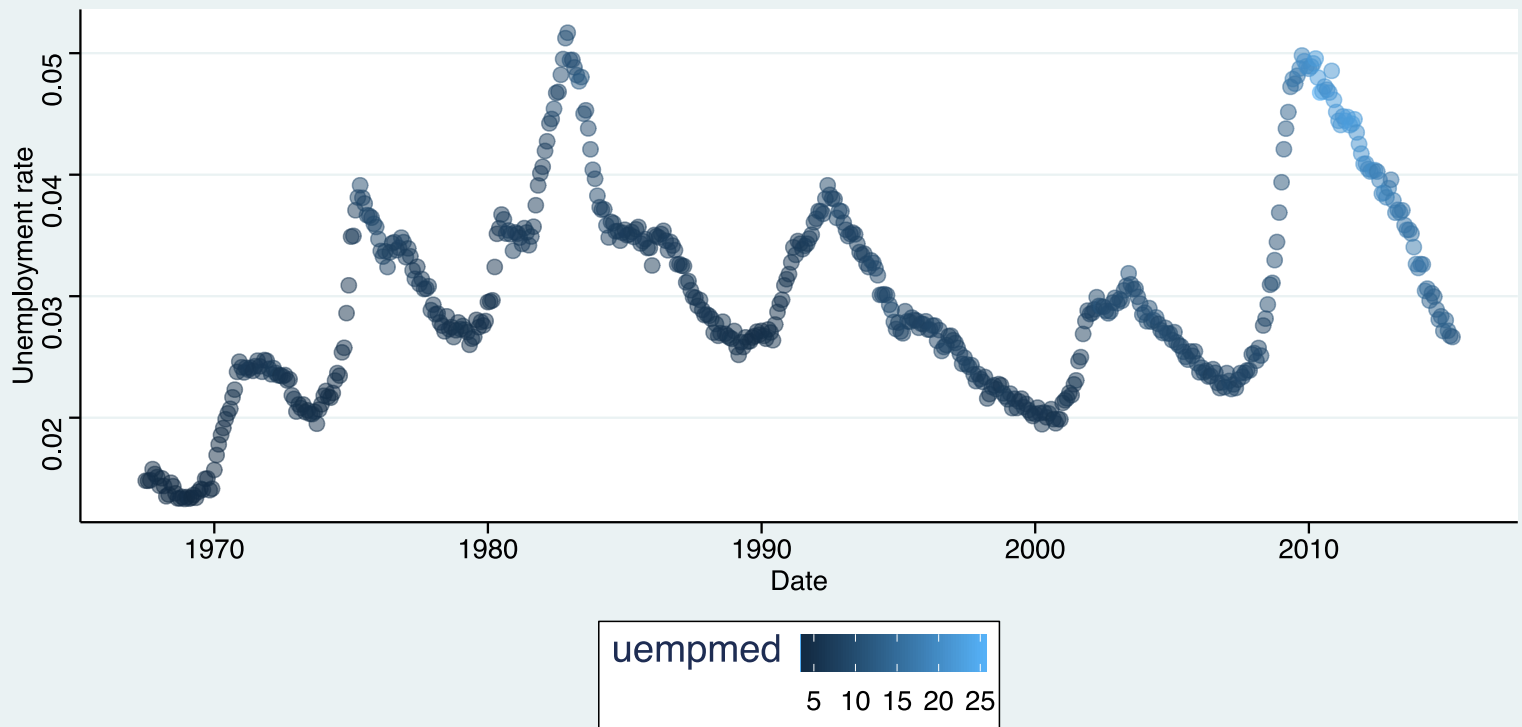
Change our theme—maybe you're a minimalist (but want slightly larger fonts)?

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed), alpha = 0.5, size = 3) +  
  theme_minimal(base_size = 14)
```



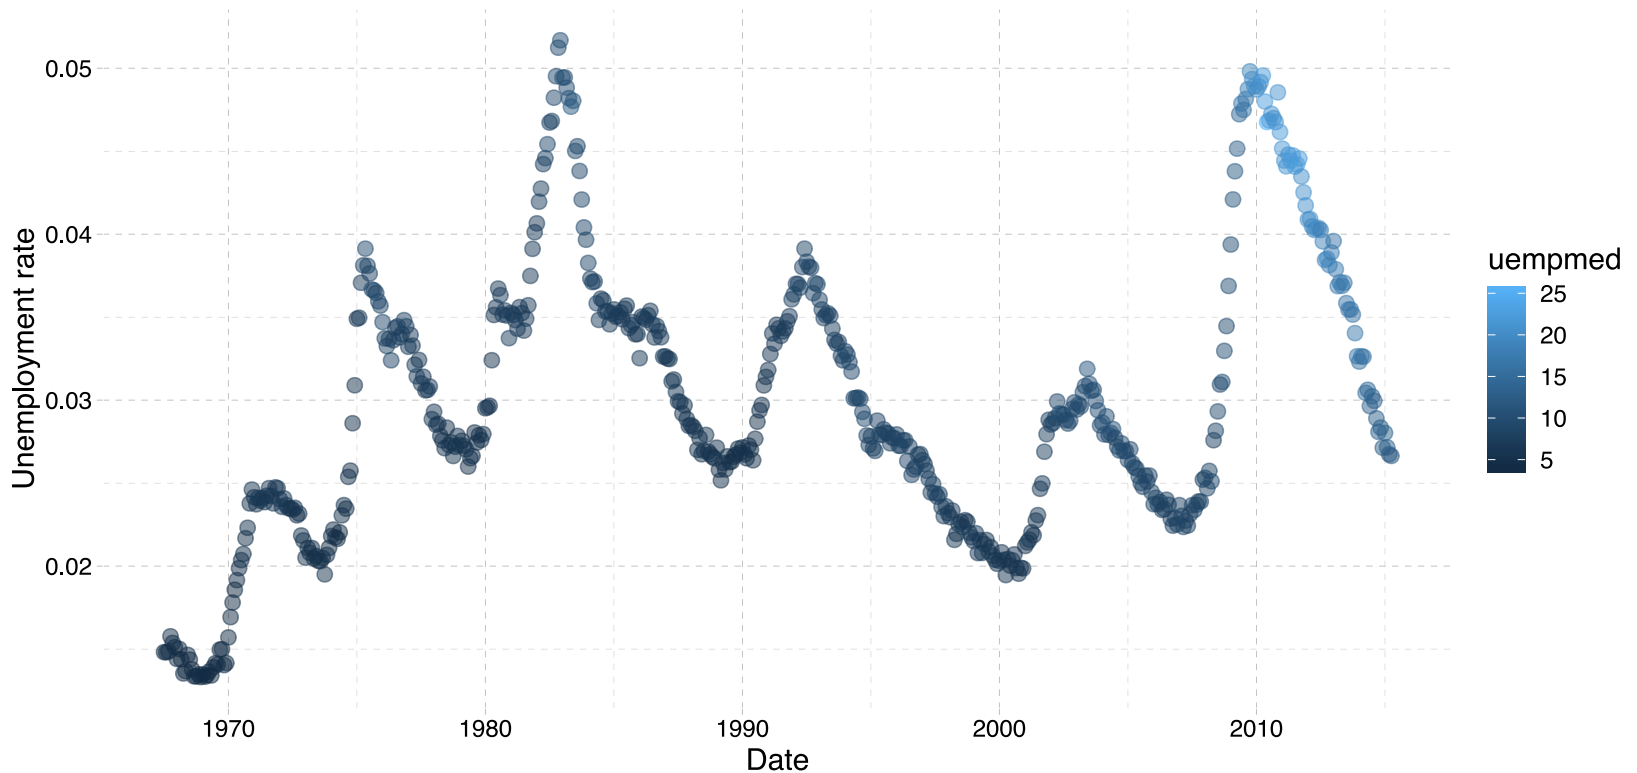
Want your figure to look like Stata made it?

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed), alpha = 0.5, size = 3) +  
  ggthemes::theme_stata(base_size = 14)
```



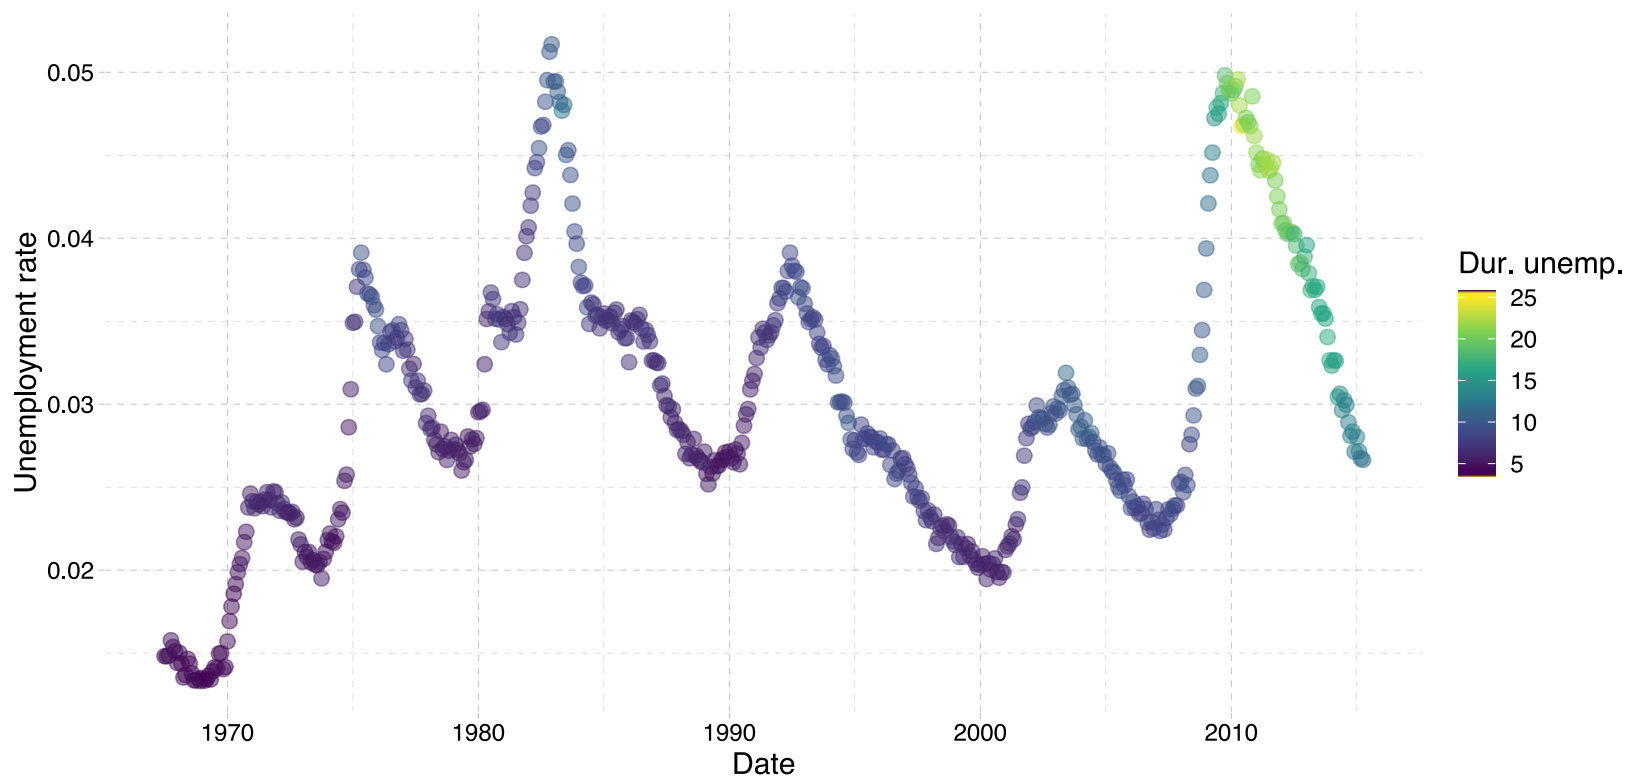
The "pander" theme from the `ggthemes` package.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed), alpha = 0.5, size = 3) +  
  ggthemes::theme_pander(base_size = 14)
```



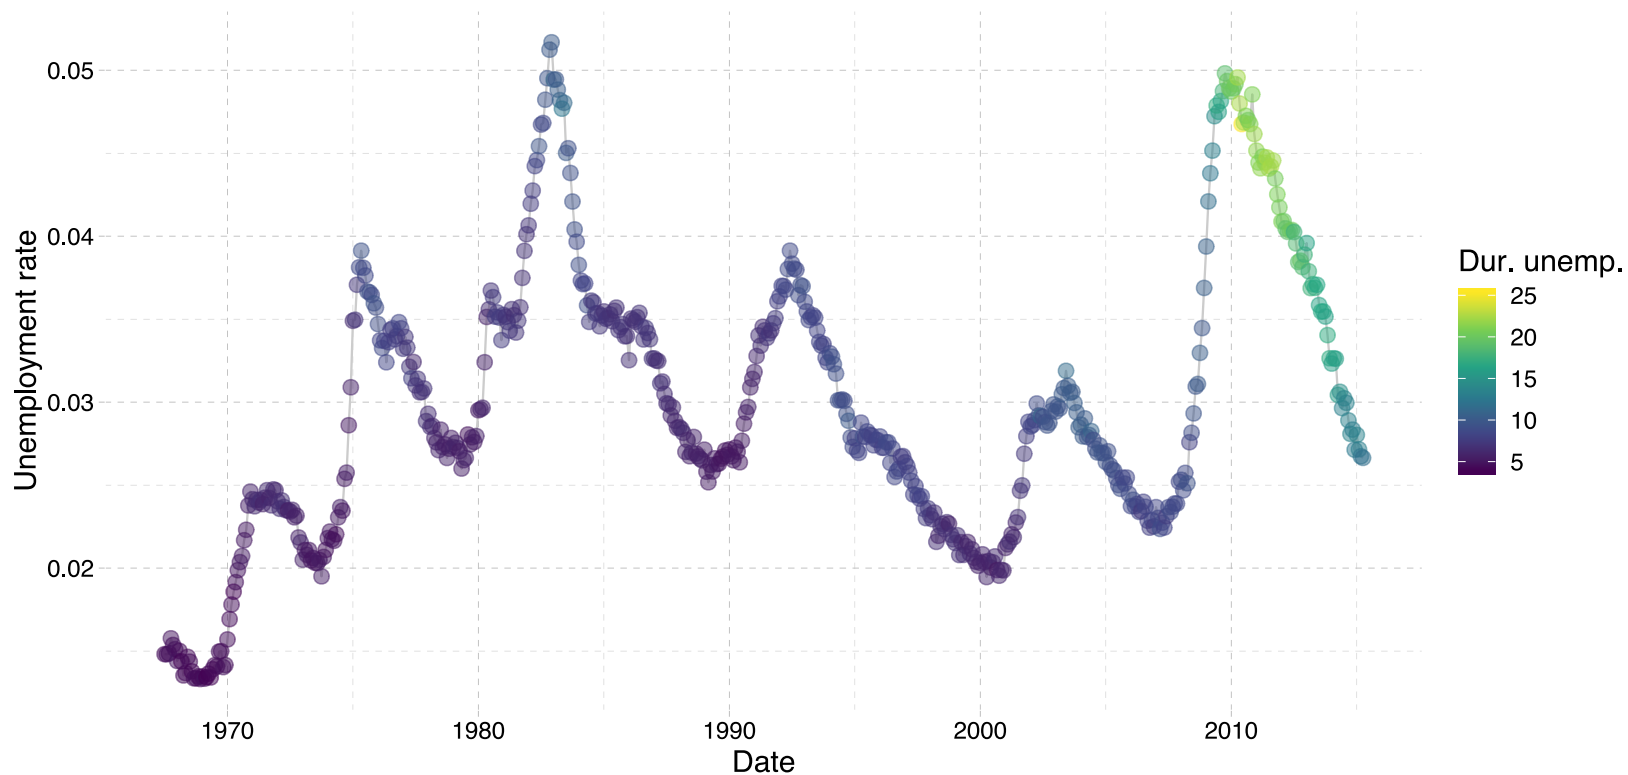
Change (and label) our color scale. Note `viridis` is the best.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed), alpha = 0.5, size = 3) +  
  ggthemes::theme_pander(base_size = 14) +  
  scale_color_viridis_c("Dur. unemp.")
```



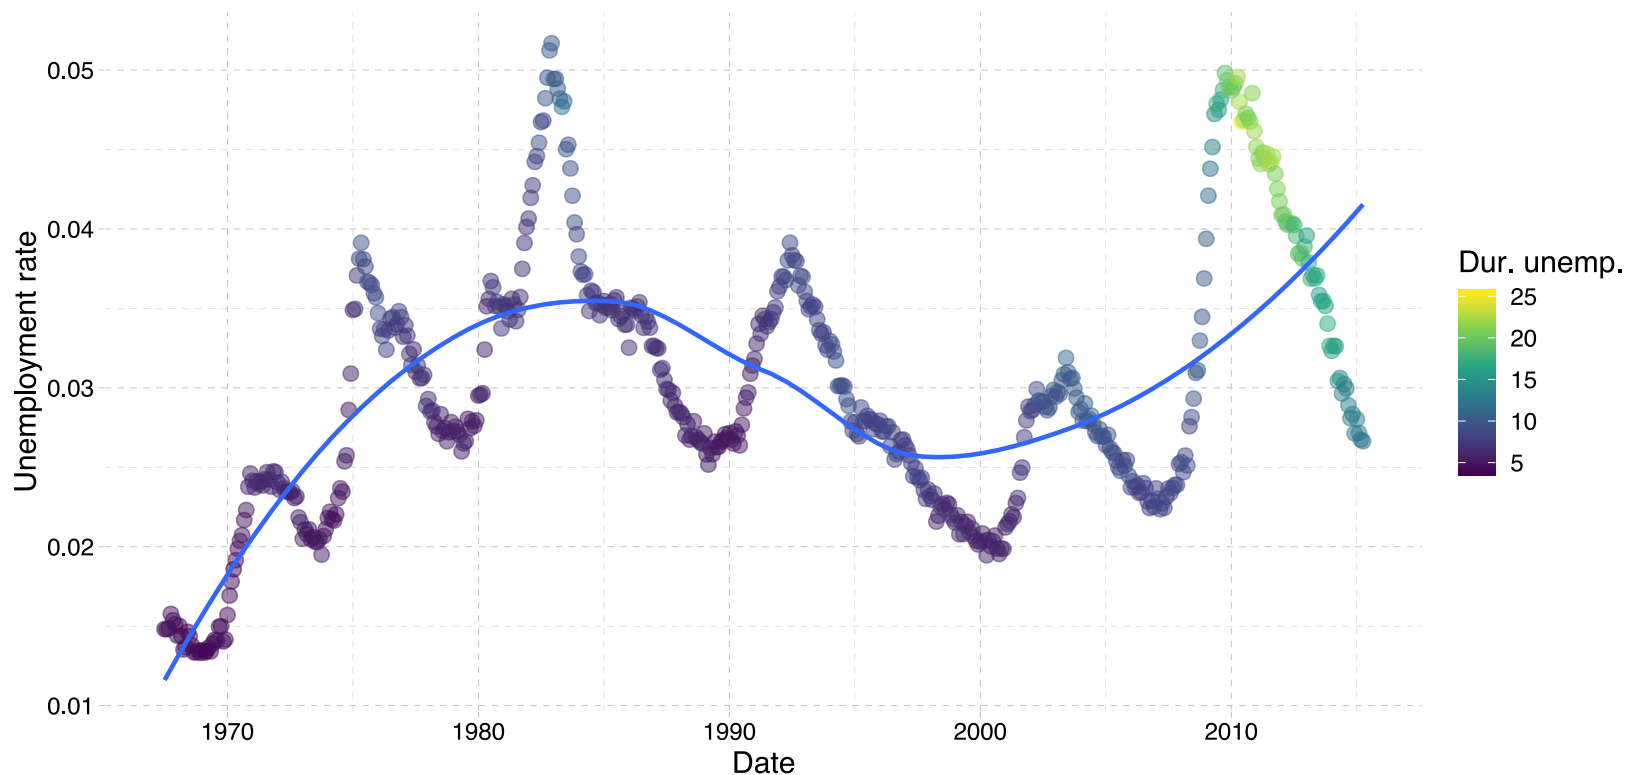
Connect the dots.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_line(color = "grey80") +  
  geom_point(aes(color = uempmed), alpha = 0.5, size = 3) +  
  ggthemes::theme_pander(base_size = 14) +  
  scale_color_viridis_c("Dur. unemp.")
```



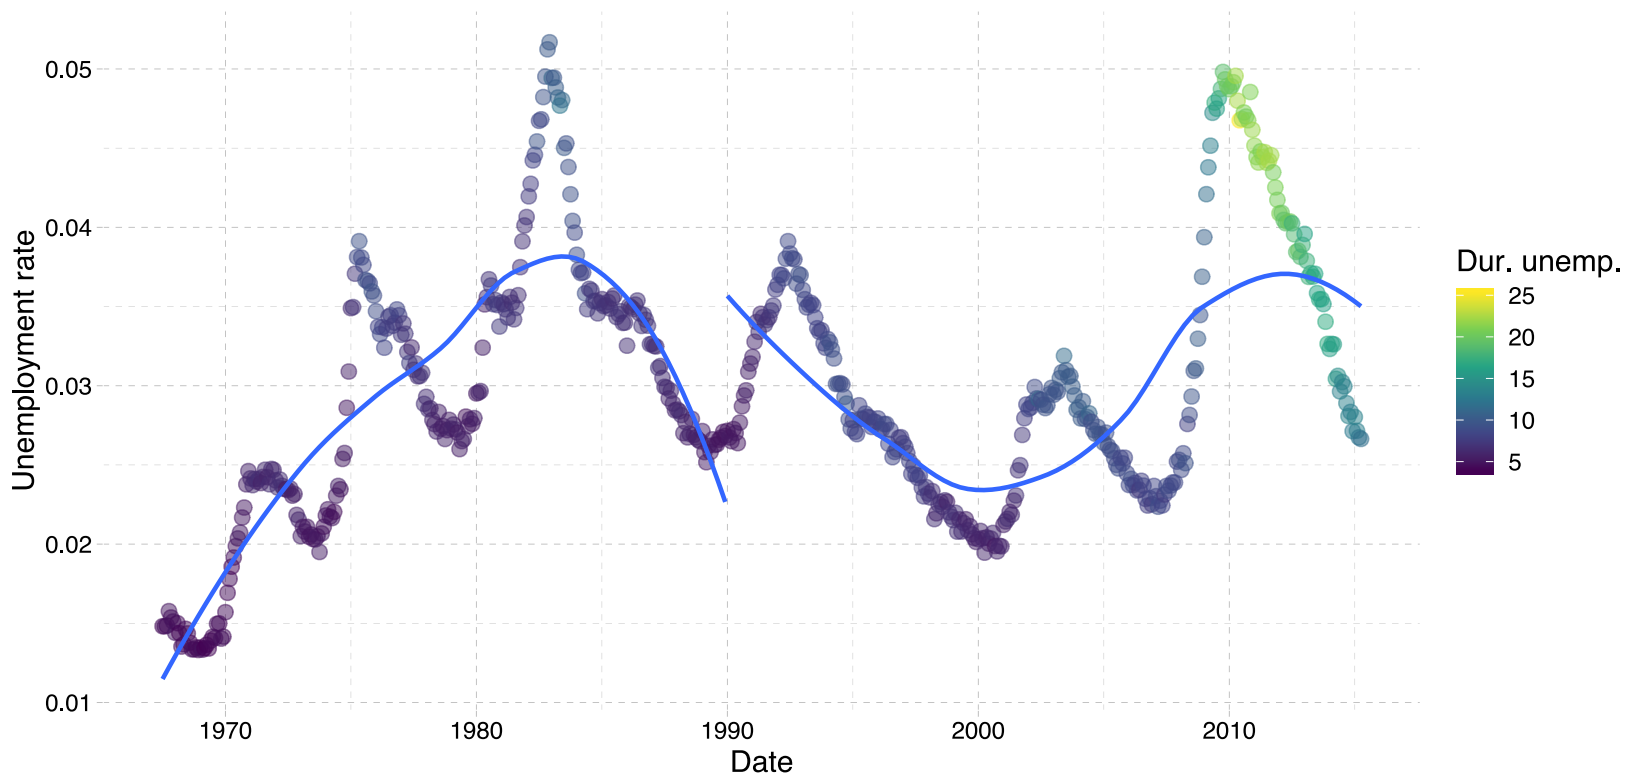
How about a smoother?

```
ggplot(data = economics, aes(x = date, y = unemploy/pop)) +  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed), alpha = 0.5, size = 3) +  
  geom_smooth(se = F) +  
  ggthemes::theme_pander(base_size = 14) +  
  scale_color_viridis_c("Dur. unemp.")
```



The `group` aesthetic separates groups.

```
ggplot(data = economics, aes(x = date, y = unemploy/pop, group = date < ymd(19900101)))  
  ylab("Unemployment rate") + xlab("Date") +  
  geom_point(aes(color = uempmed), alpha = 0.5, size = 3) +  
  geom_smooth(se = F) +  
  ggthemes::theme_pander(base_size = 14) +  
  scale_color_viridis_c("Dur. unemp.")
```

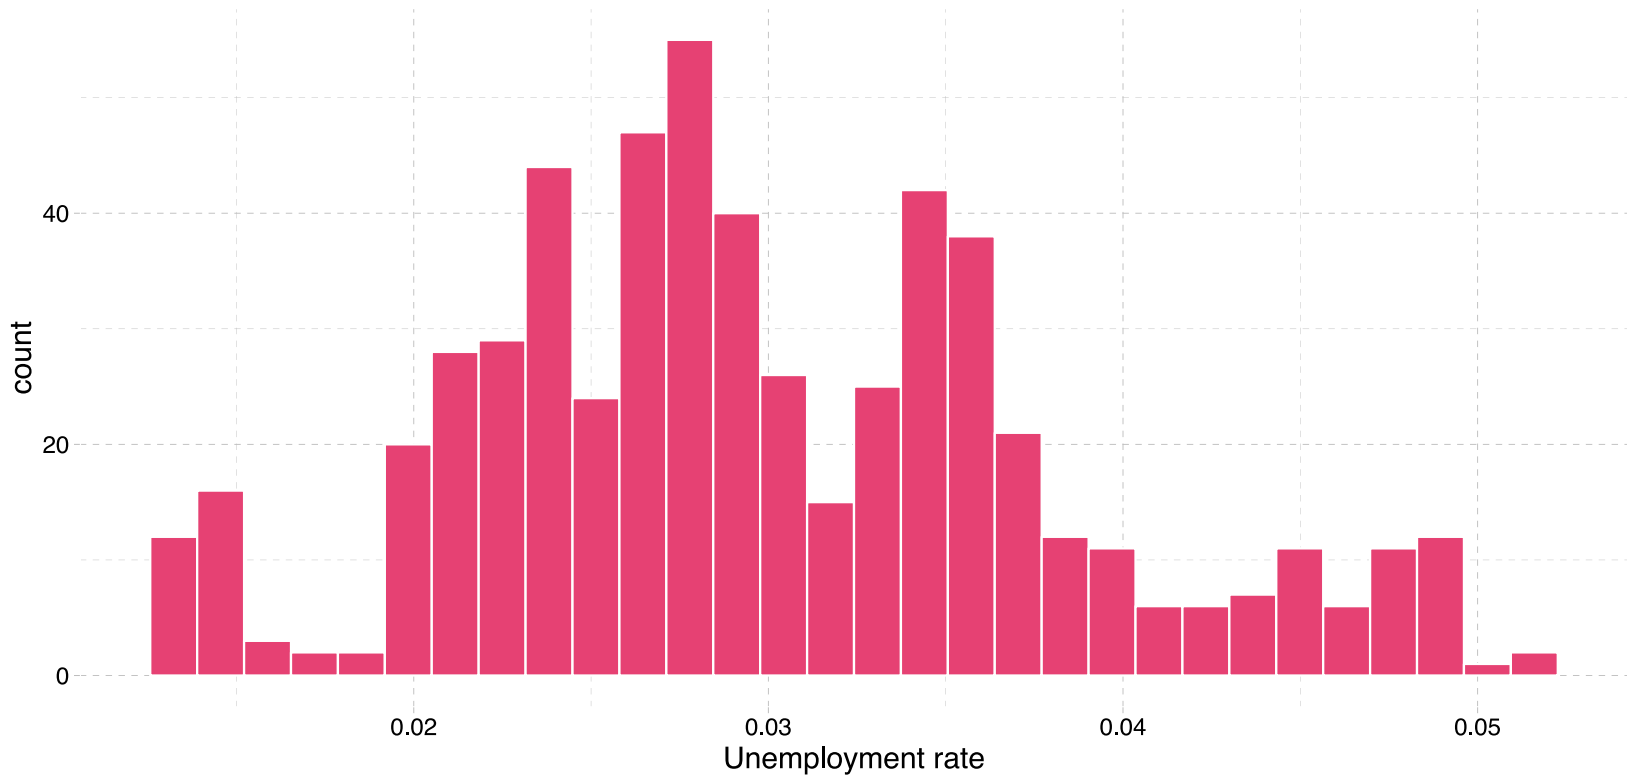


Note The `ymd()` function comes from the `lubridate` package.

`ggplot2` knows histograms.

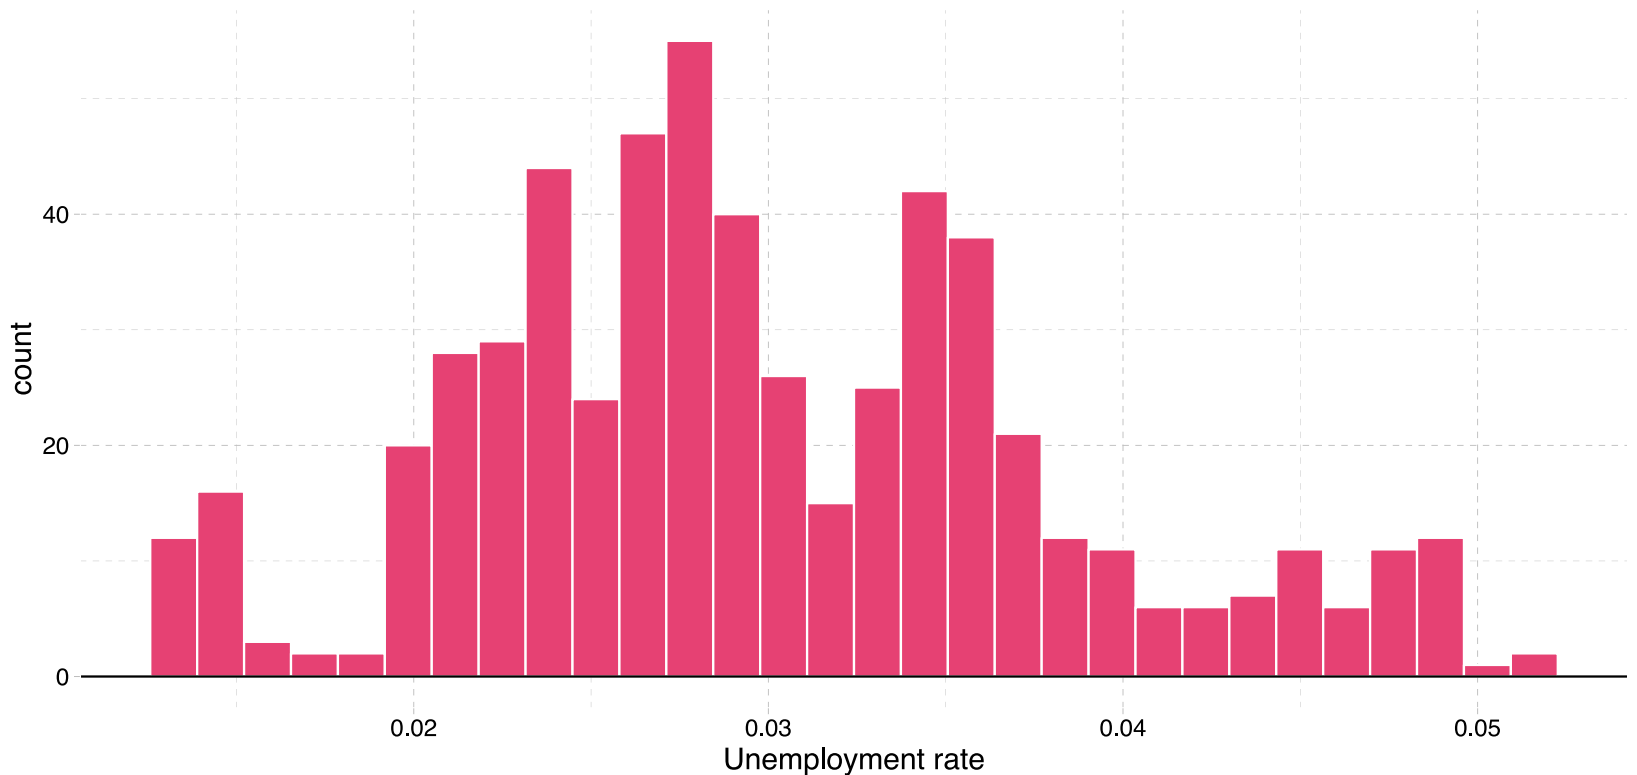
A histogram.

```
ggplot(data = economics, aes(x = unemploy/pop)) +  
  xlab("Unemployment rate") +  
  geom_histogram(color = "white", fill = "#e64173") +  
  ggthemes::theme_pander(base_size = 14)
```



Add a horizontal line where count = 0.

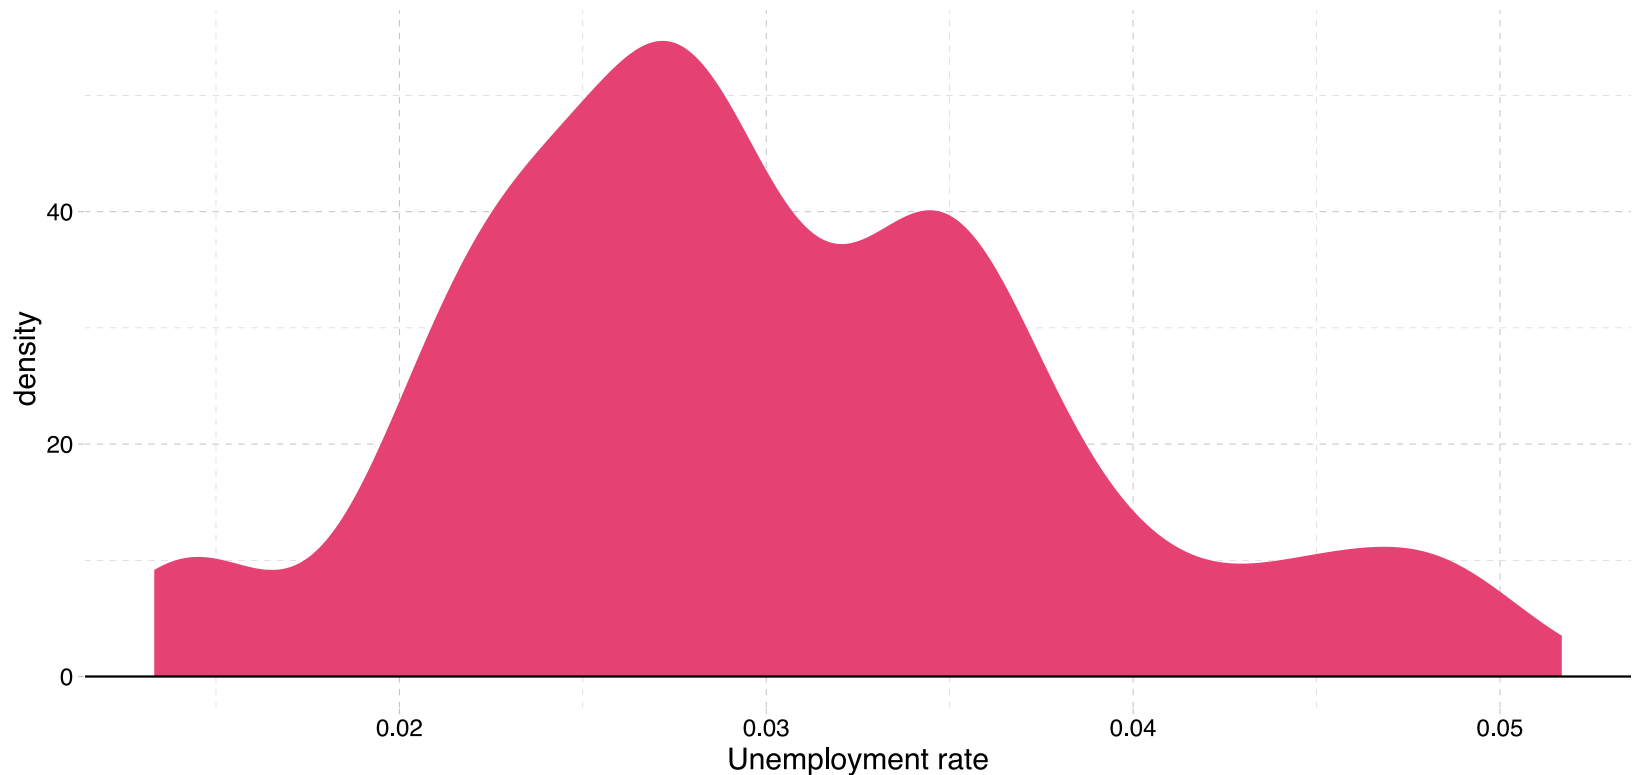
```
ggplot(data = economics, aes(x = unemploy/pop)) +  
  xlab("Unemployment rate") +  
  geom_histogram(color = "white", fill = "#e64173") +  
  geom_hline(yintercept = 0) +  
  ggthemes::theme_pander(base_size = 14)
```



`ggplot2` knows densities.

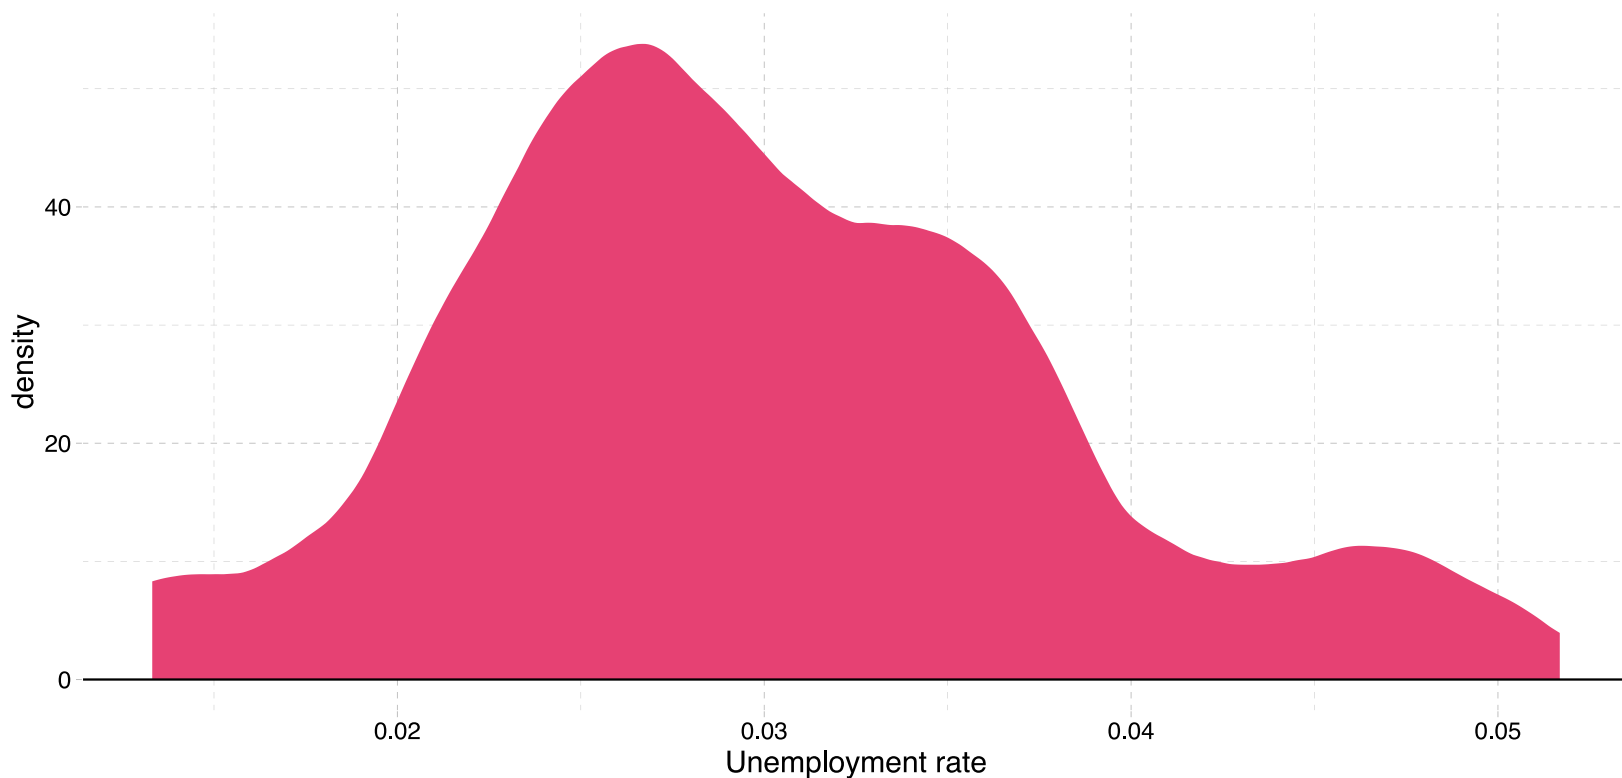
A density plot.

```
ggplot(data = economics, aes(x = unemploy/pop)) +  
  xlab("Unemployment rate") +  
  geom_density(color = NA, fill = "#e64173") +  
  geom_hline(yintercept = 0) +  
  ggthemes::theme_pander(base_size = 14)
```



Now with Epanechnikov kernel!

```
ggplot(data = economics, aes(x = unemploy/pop)) +  
  xlab("Unemployment rate") +  
  geom_density(kernel = "epanechnikov", color = NA, fill = "#e64173") +  
  geom_hline(yintercept = 0) +  
  ggthemes::theme_pander(base_size = 14)
```

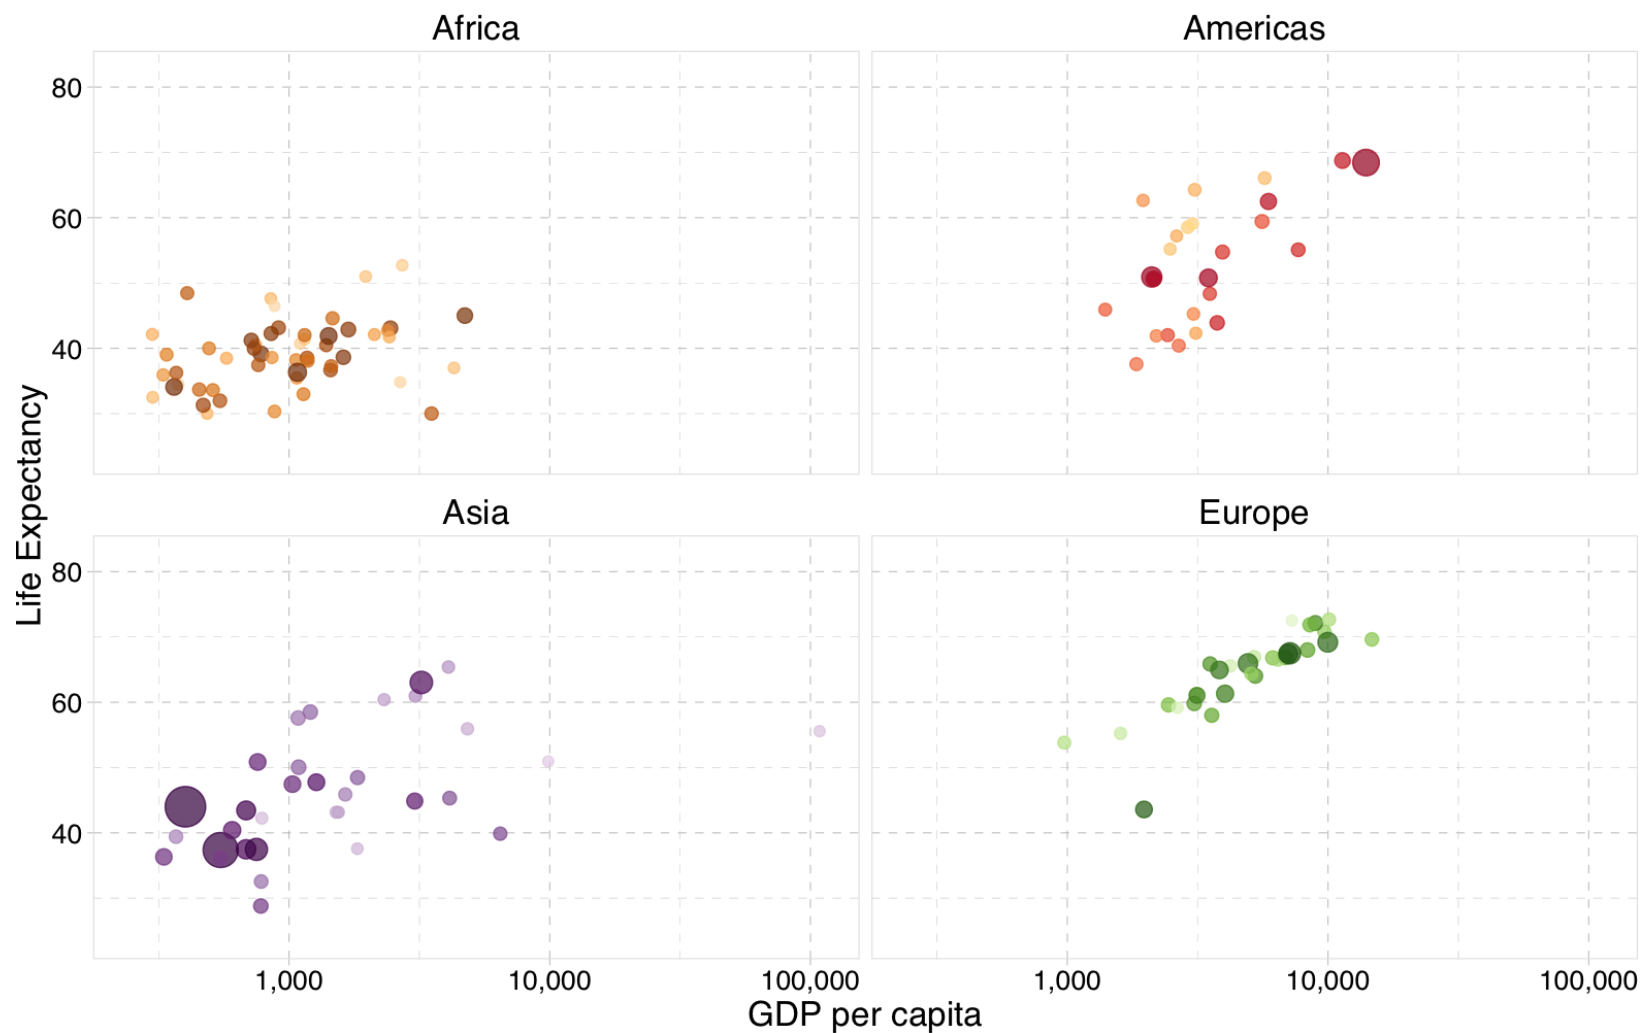


`ggplot2` itself is incredibly flexible/powerful.

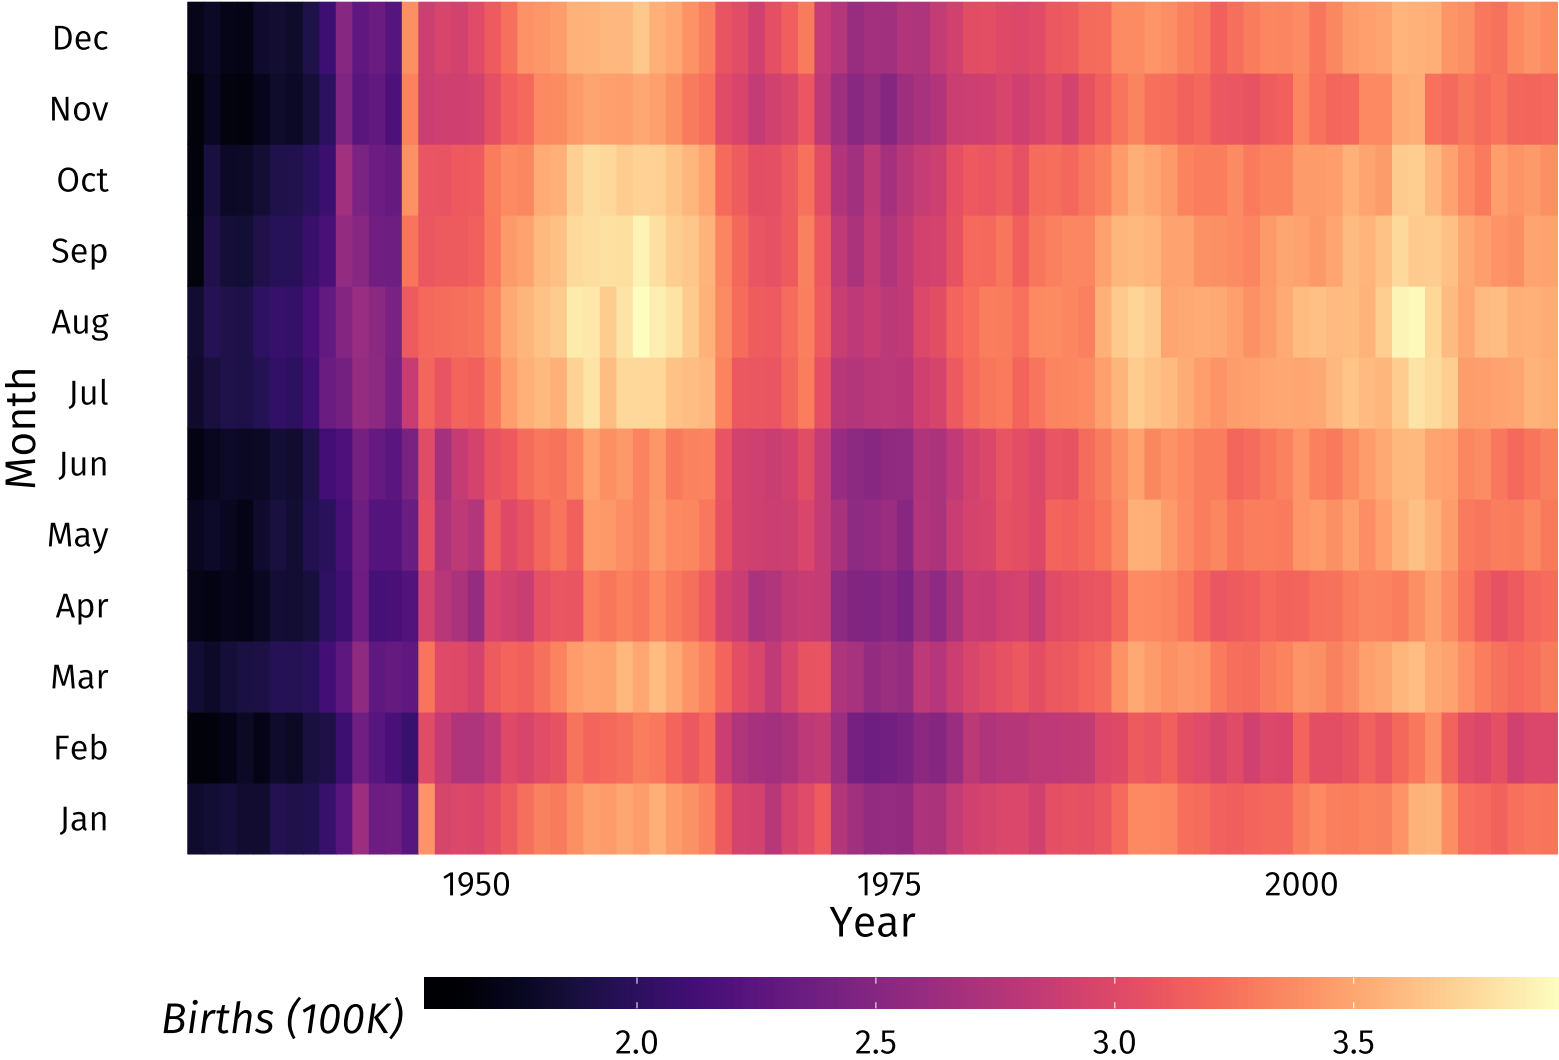
But there are **even more packages** that extend its power—e.g., `ggthemes`, `gganimate`, `cowplot`, `ggmap`, `ggExtra`, and (of course) `viridis`.

Gapminder meets `gganimate`

Year: 1952



US births by month since 1933



Resources

There's always more

ggplot2

1. RStudio's [cheat sheet for ggplot2](#).
2. [ggplot2 reference index](#)
3. The [tidyverse page](#) on [ggplot2](#).
4. Hadley Wickham's on *Data visualization* in his data science book.

Table of contents

Plotting

1. `plot()`
 - Description
 - Examples
 - Layering plots
2. `hist()`
3. `ggplot2`
 - Intro
 - `ggplot()`
 - Layers
 - Building a plot
 - Histogram
 - Density
 - More
4. More resources